

# Metrics & Monitoring

---

Konstantin Knauf, Solutions Architect

# Metrics & Monitoring

## Agenda

- Flink's Metrics System *"How"*
  - Metrics
  - MetricsReporter
- Key Metrics for Continuous Monitoring *"What"*
  - Health
  - Throughput & Progress
  - Latency
- Key Metrics for Troubleshooting *"What else"*



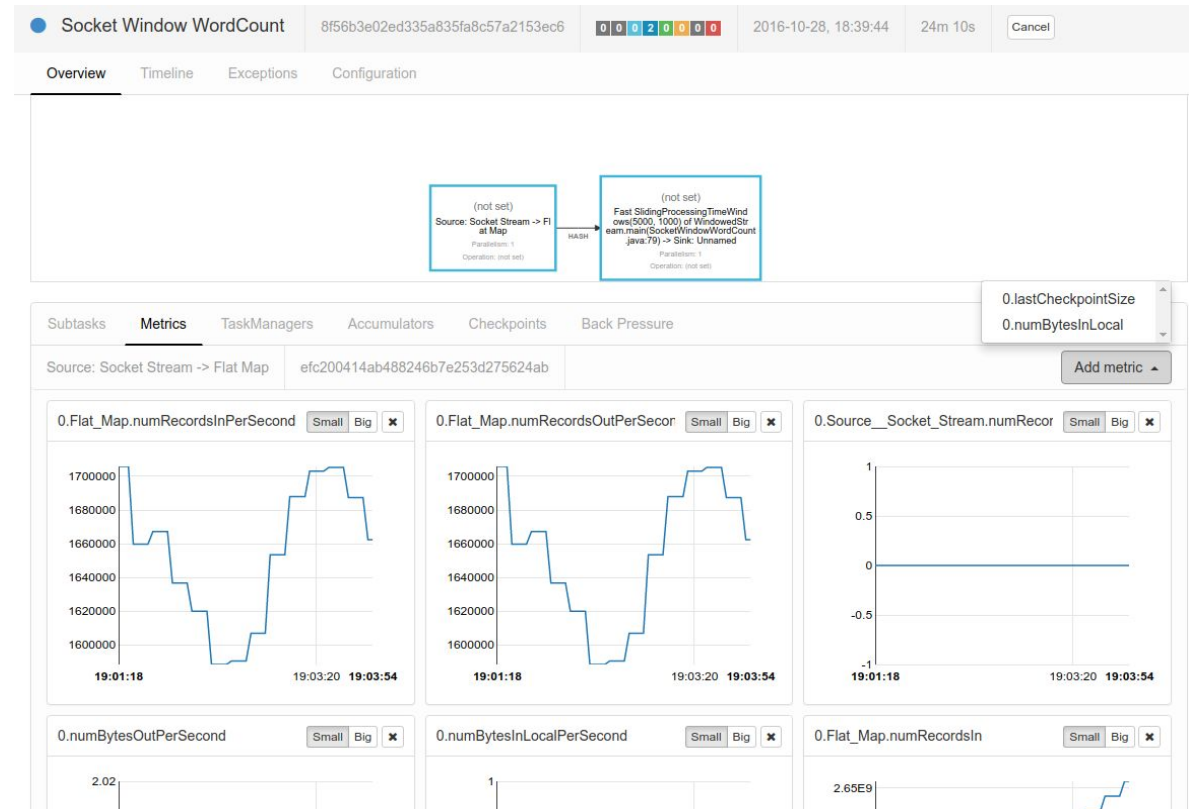
# Flink's Metrics System

---



# Metrics

- <identifier, measurement>
- Types
  - Counter
  - Meter (rate)
  - Histogram
  - Gauge (arbitrary value)



# Example

```
public static class MyMap extends RichMapFunction<String, String> {  
    private Counter count;  
  
    @Override  
    public void open(Configuration config) {  
        count = getRuntimeContext()  
            .getMetricGroup()  
            .counter("numRecordsIn");  
    }  
  
    @Override  
    public String map(String input) {  
        count.inc();  
        // return something  
    }  
}
```



# Metrics

## Scopes

- metrics scope to different levels of a Flink deployment
- the keys to attach to metrics in a certain scope can be configured
  - metrics.scope.jm: <host>.jobmanager
  - metrics.scope.task:  
<host>.taskmanager.<tm\_id>.<job\_name>.<task\_name>.<subtask\_index>
- Checkout  
<https://ci.apache.org/projects/flink/flink-docs-release-1.7/monitoring/metrics.html#scope> for details



# Accessing Metrics

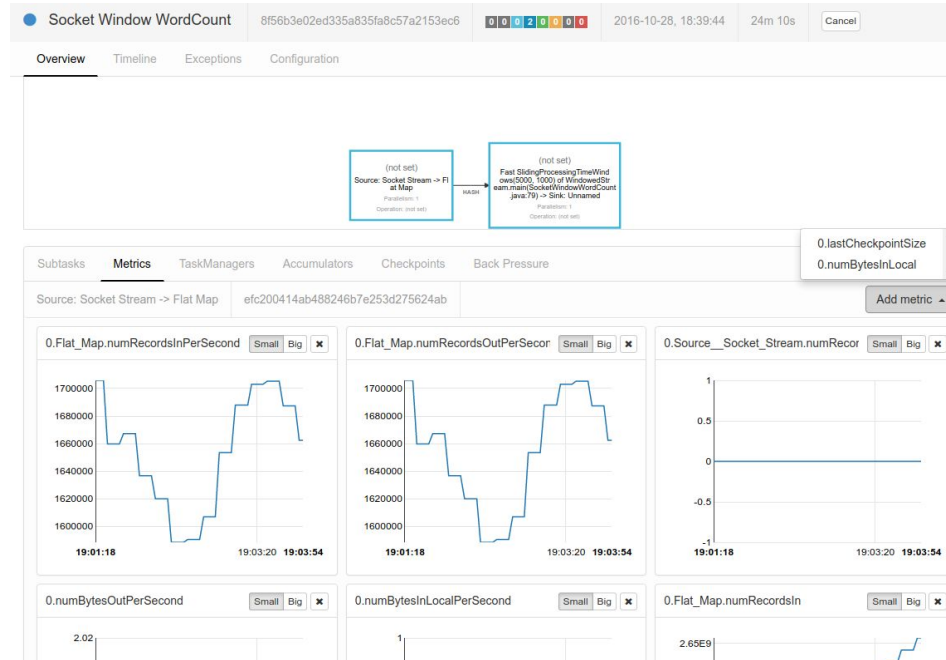
- WebUI → TaskMetrics
- REST API
- **MetricsReporters**

/jobs/<id>/metrics

/jobs/<id>/checkpoints

/jobs/<id>/vertices/<id>/metrics?get=0.numRecordsOutPerSecond

/taskmanagers/<id>/metrics?get=<metric>



# Accessing Metrics

## Metrics Reporters

- Datadog
- Ganglia
- Graphite
- JMX



- Prometheus
- StatsD
- SLF4J
- InfluxDB



Or write your own...



# Accessing Metrics

## A Simple Log4jReporter

```
public static class Log4JReporter implements MetricReporter, Scheduled {  
    private static final Logger LOG = LoggerFactory.getLogger(Log4jReporter.class);  
  
    private final Map<Counter, String> counters = new ConcurrentHashMap<>();  
  
    public void notifyOfAddedMetric(Metric metric, String metricName, MetricGroup group) {  
        if (metric instanceof Counter) {  
            counters.put((Counter) metric, group.getMetricIdentifier(metricName));  
        }  
    }  
  
    public void notifyOfRemovedMetric(Metric metric, String metricName, MetricGroup group) {  
        if (metric instanceof Counter) {  
            counters.remove(metric);  
        }  
    }  
  
    public void report() {  
        for (Map.Entry<Counter, String> metric : counters.entrySet()) {  
            LOG.info(metric.getValue() + ": " + metric.getKey());  
        }  
    }  
}
```



# Key Metrics for Continuous Monitoring

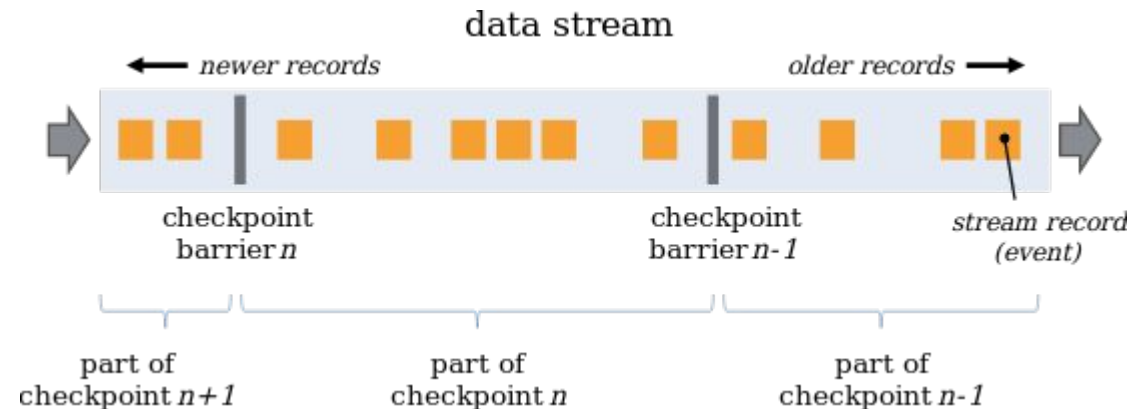
---



# Key Metrics

## General Health

- Is “RUNNING”?
  - uptime
  - fullRestarts
- Checkpointing Consistently?
  - numberOfCompletedCheckpoints
  - numberOfFailedCheckpoints
  - lastCheckpointSize



# Key Metrics

## Throughput & Progress

- Task & Operator Level Throughput
  - `numRecords(In|Out)PerSecond`
  - `numRecords(In|Out)`
- Progress & Event-Time Lag
  - `currentOutputWatermark`
- Keeping Up
  - (Kafka) `records-lag-max`
  - (Kinesis) `millisBehindLatest`



# Key Metrics

## Latency

- Add timestamp to events at multiple stages, e.g.
  - event creation
  - ingestion
  - publishing
- custom metrics for reporting these metrics



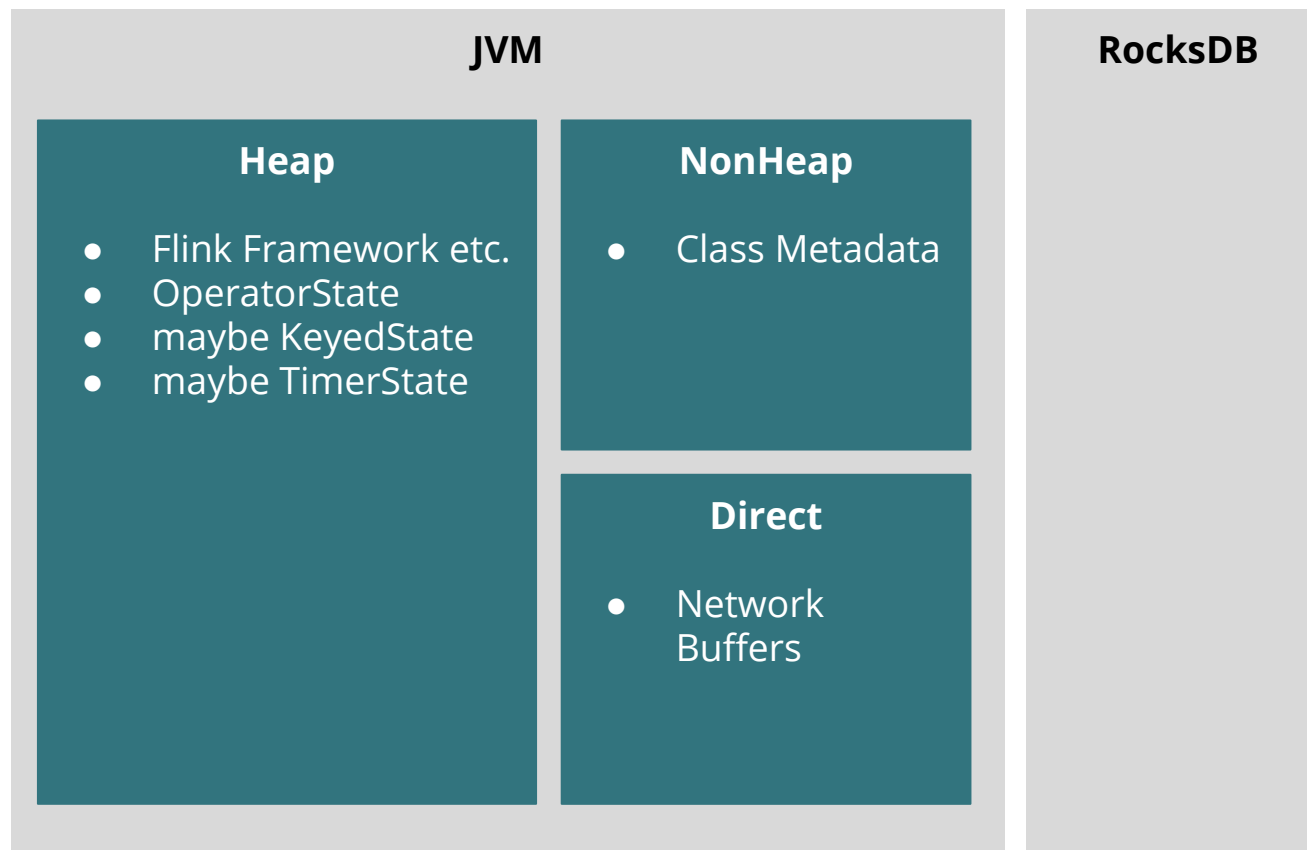
# Key Metrics for Troubleshooting



# JVM Metrics

## Memory

- `Status.JVM.Memory.`
  - `NonHeap.Committed`
  - `Heap.Used`
  - `Heap.Committed`
  - `Direct.MemoryUsed`
  - `Mapped.MemoryUsed`
  - `G1 Young Generation.Time`
  - `G1 Old Generation.Time`



# JVM Metrics

## CPU

- Metrics
    - `Status.JVM.CPU.Load`
    - `Status.JVM.CPU.Time`
  - Leave some slack for catch-up scenarios (& RocksDB)
- Note:** 0.021 = 100% load for a Taskmanager container with 1 CPU on a 48 core machine.






# Troubleshooting Latency

## Latency Tracking


- For each operator-subtask a latency histogram is exposed
- Enabled via `metrics.latency.interval`
- scoped to job
- `latency.source_id.<source_id>.operator_id.<operator_id>.operator_subtask_index.<subtask_index>.`



# Troubleshooting Latency

## Latency Tracking

- For each operator-subtask a latency histogram is exposed
- Enabled via `metrics.latency.interval`




# Troubleshooting Backpressure

## Latency Tracking

**metrics.latency.granularity: single**

- Per Subtask
  - Latency histogram for both sources
- Overall
  - $4 (P * \#Operators)$




# Troubleshooting Backpressure

## Latency Tracking

**metrics.latency.granularity: operator**

- Per Subtask
  - Latency histogram for Source A
  - Latency histogram for Source B
- Overall
  - 8 histograms ( $P * \text{\#Sources} * \text{\#Operators}$ )




# Troubleshooting Backpressure

## Latency Tracking

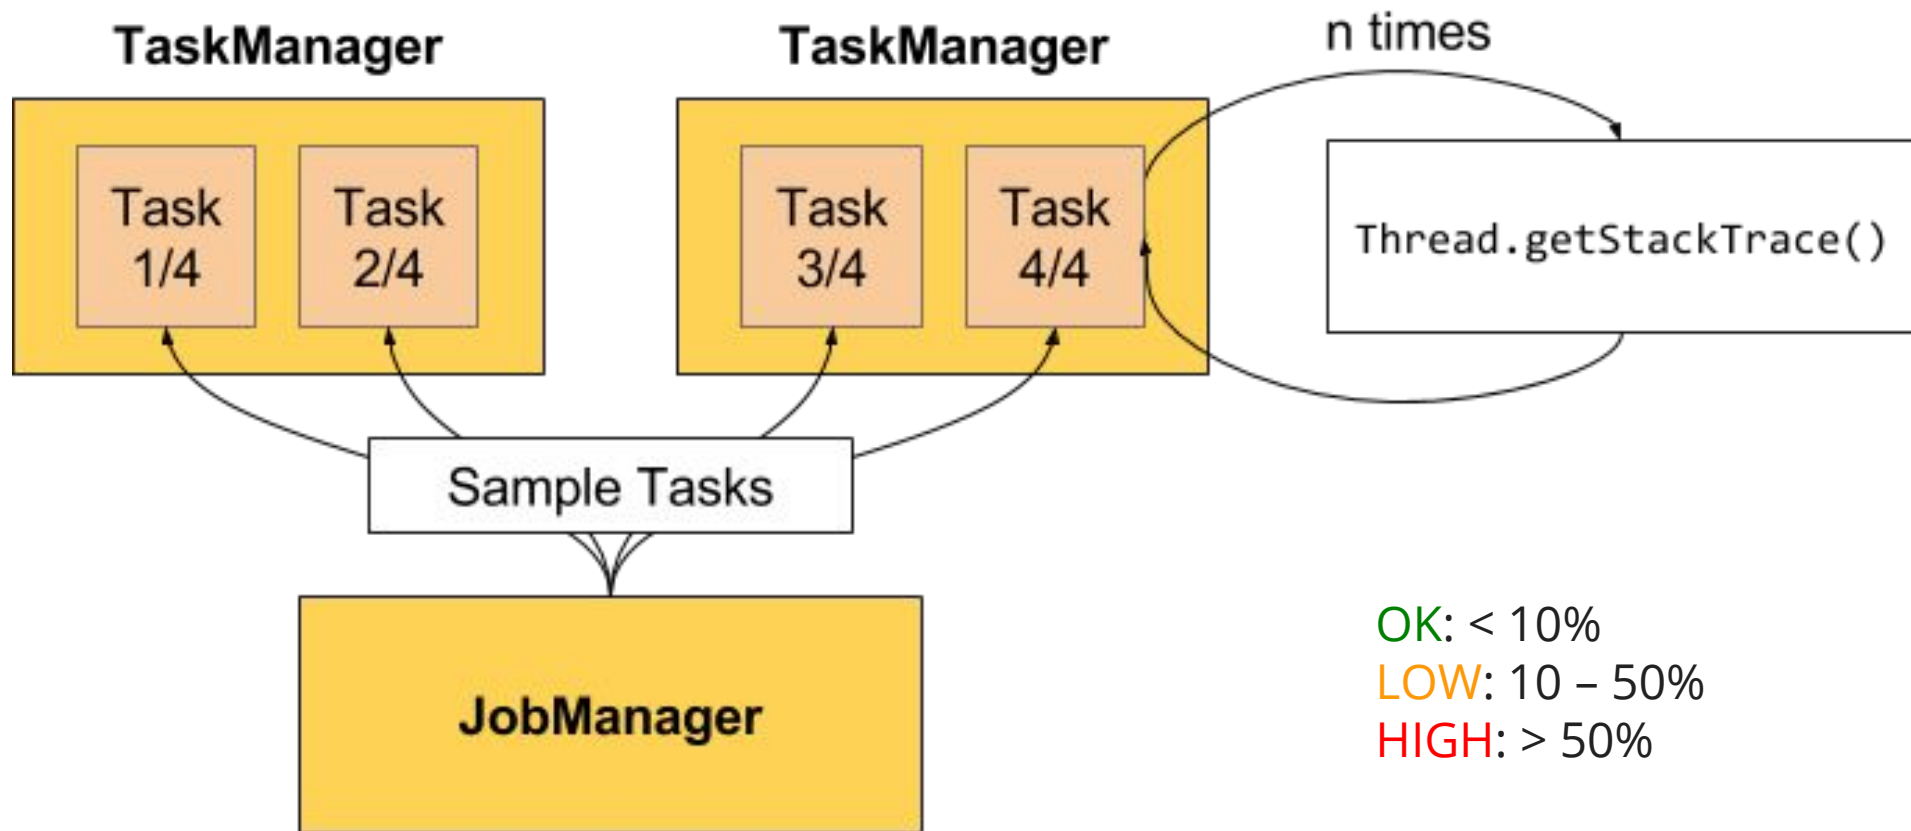
**metrics.latency.granularity: subtask**

- Per Subtask
  - Latency histogram for Source A (1/2)
  - Latency histogram for Source A (2/2)
  - Latency histogram for Source B (1/2)
  - Latency histogram for Source B (2/2)
- Overall
  - 16 histogram ( $P^2 * \text{\#Sources} * \text{\#Operators}$ )



# Troubleshooting Backpressure

## Backpressure Monitor





localhost:8081/#/jobs/156c3b5d

Flink Streaming Job

00000020

37s

Cancel

Plan

Timeline

Exceptions

Properties

Configuration

Source: Custom Source

RUNNING

Measurement

Back Pressure Status

10s ago

OK

Hide subtasks ^

^

Subtask	Ratio	Status
1	0.03	OK
2	0.03	OK
3	0.03	OK
4	0.04	OK

Flat Map -> Sink: Unnamed

RUNNING





localhost:8081/#/jobs/156c3b5d

Flink Streaming Job

00000020

1m 37s

Cancel

Plan

Timeline

Exceptions

Properties

Configuration

Source: Custom Source

RUNNING

Measurement

Back Pressure Status

2s ago

HIGH

Hide subtasks ^

^

Subtask	Ratio	Status
1	1	HIGH
2	1	HIGH
3	1	HIGH
4	1	HIGH

Flat Map -> Sink: Unnamed


RUNNING



# Troubleshooting Backpressure

## Asymmetric Backpressure

- situation where backpressure only occurs in one channel
- hard to detect, but can lead to checkpoint timeouts
- Metrics
  - `inputQueueLength`
  - `outputQueueLength`





# ververica

---

[konstantin@ververica.com](mailto:konstantin@ververica.com)

[www.ververica.com](http://www.ververica.com)

[@VervericaData](https://twitter.com/VervericaData)