

Troubleshooting & Operations

Nico Kruber, Solutions Architect & Apache Flink committer
Alexander Fedulov, Solutions Architect



Agenda

Morning

- 10:00am - 11:30am: Getting Started
 - Flink's Distributed Architecture "Recap"
 - Getting started with the hands-on exercises
- **Coffee Break**
- 11:45am - 1:15pm: Event-Time & Latency
 - Metrics & Monitoring
 - Watermarking
 - Latency



Agenda

Afternoon

- 2:15pm - 3:45pm: Performance Tuning
 - Serialization
 - Throughput Optimization
- **Coffee Break**
- 4:00pm - 5:30pm: State, Memory, & Serialization
 - State & State Backends
 - Object-Reuse & Serialization (2)

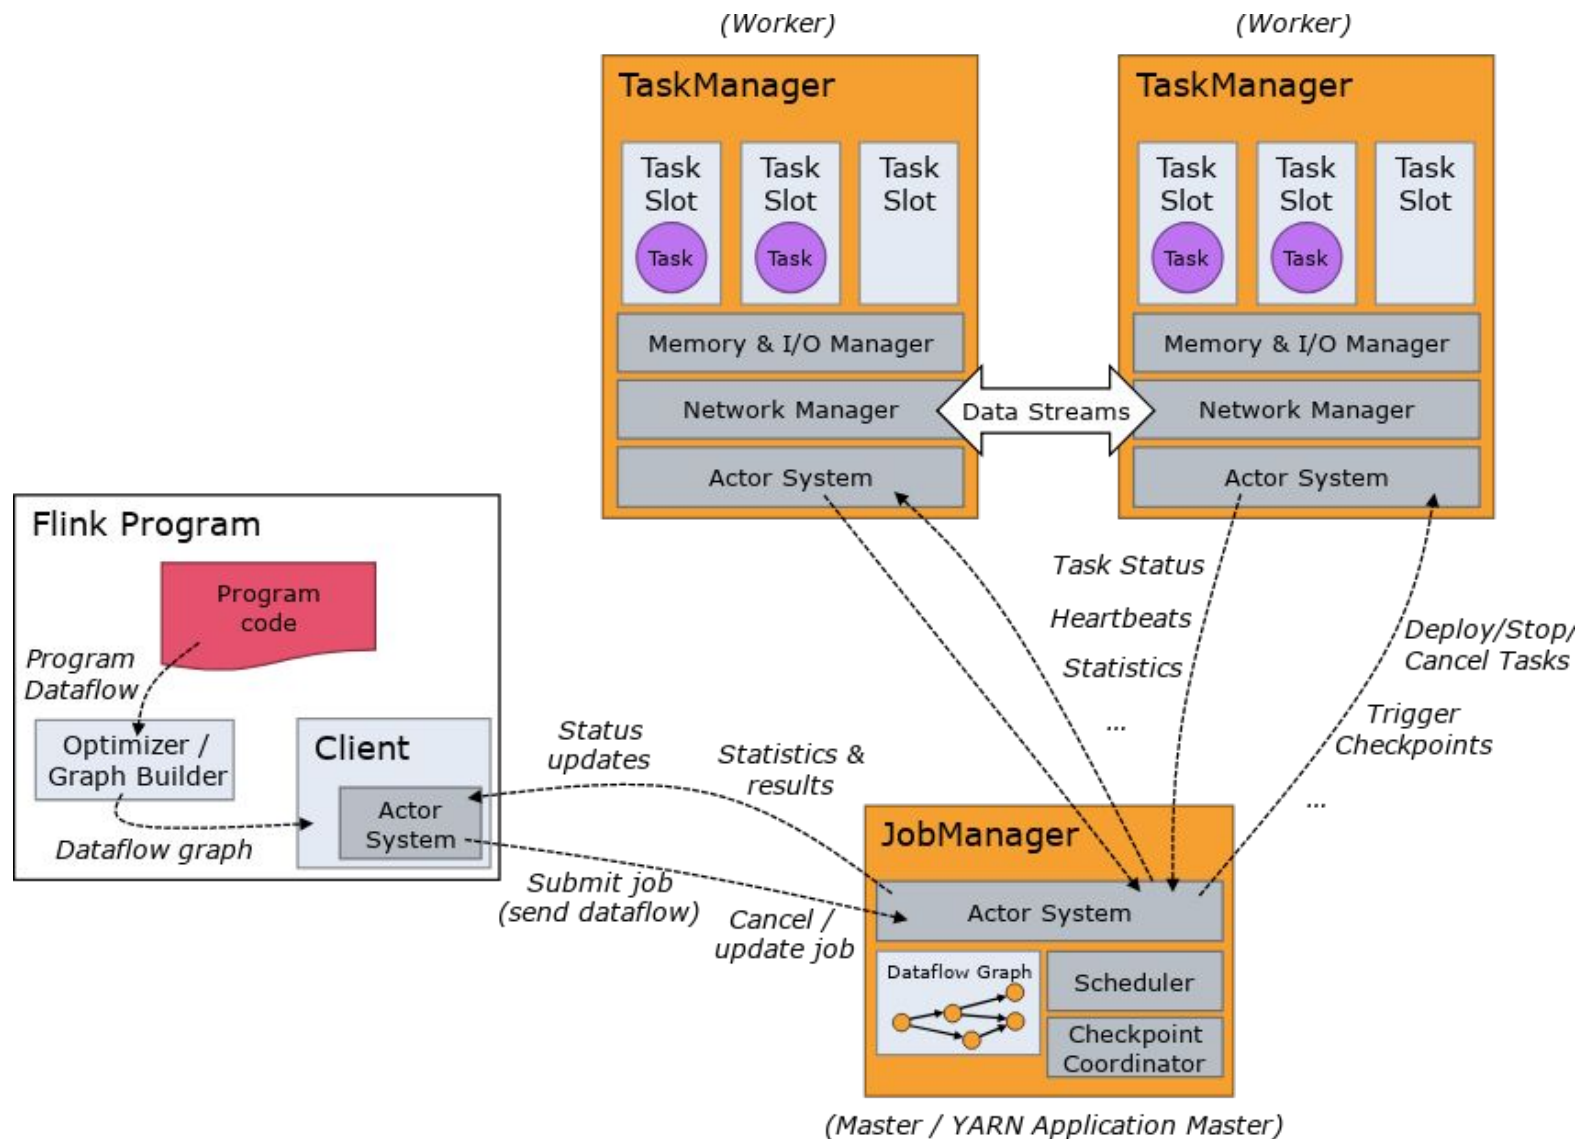


Flink's Distributed Architecture

Nico Kruber, Solutions Architect & Apache Flink committer



Cluster Components



Job Definition

```
DataStream<String> lines = env.addSource(  
    new FlinkKafkaConsumer<>(...));
```

} Source

```
DataStream<Event> events = lines.map((line) -> parse(line));
```


} Transformation

```
DataStream<Statistics> stats = events  
    .keyBy("id")  
    .timeWindow(Time.seconds(10))  
    .apply(new MyWindowAggregationFunction());
```

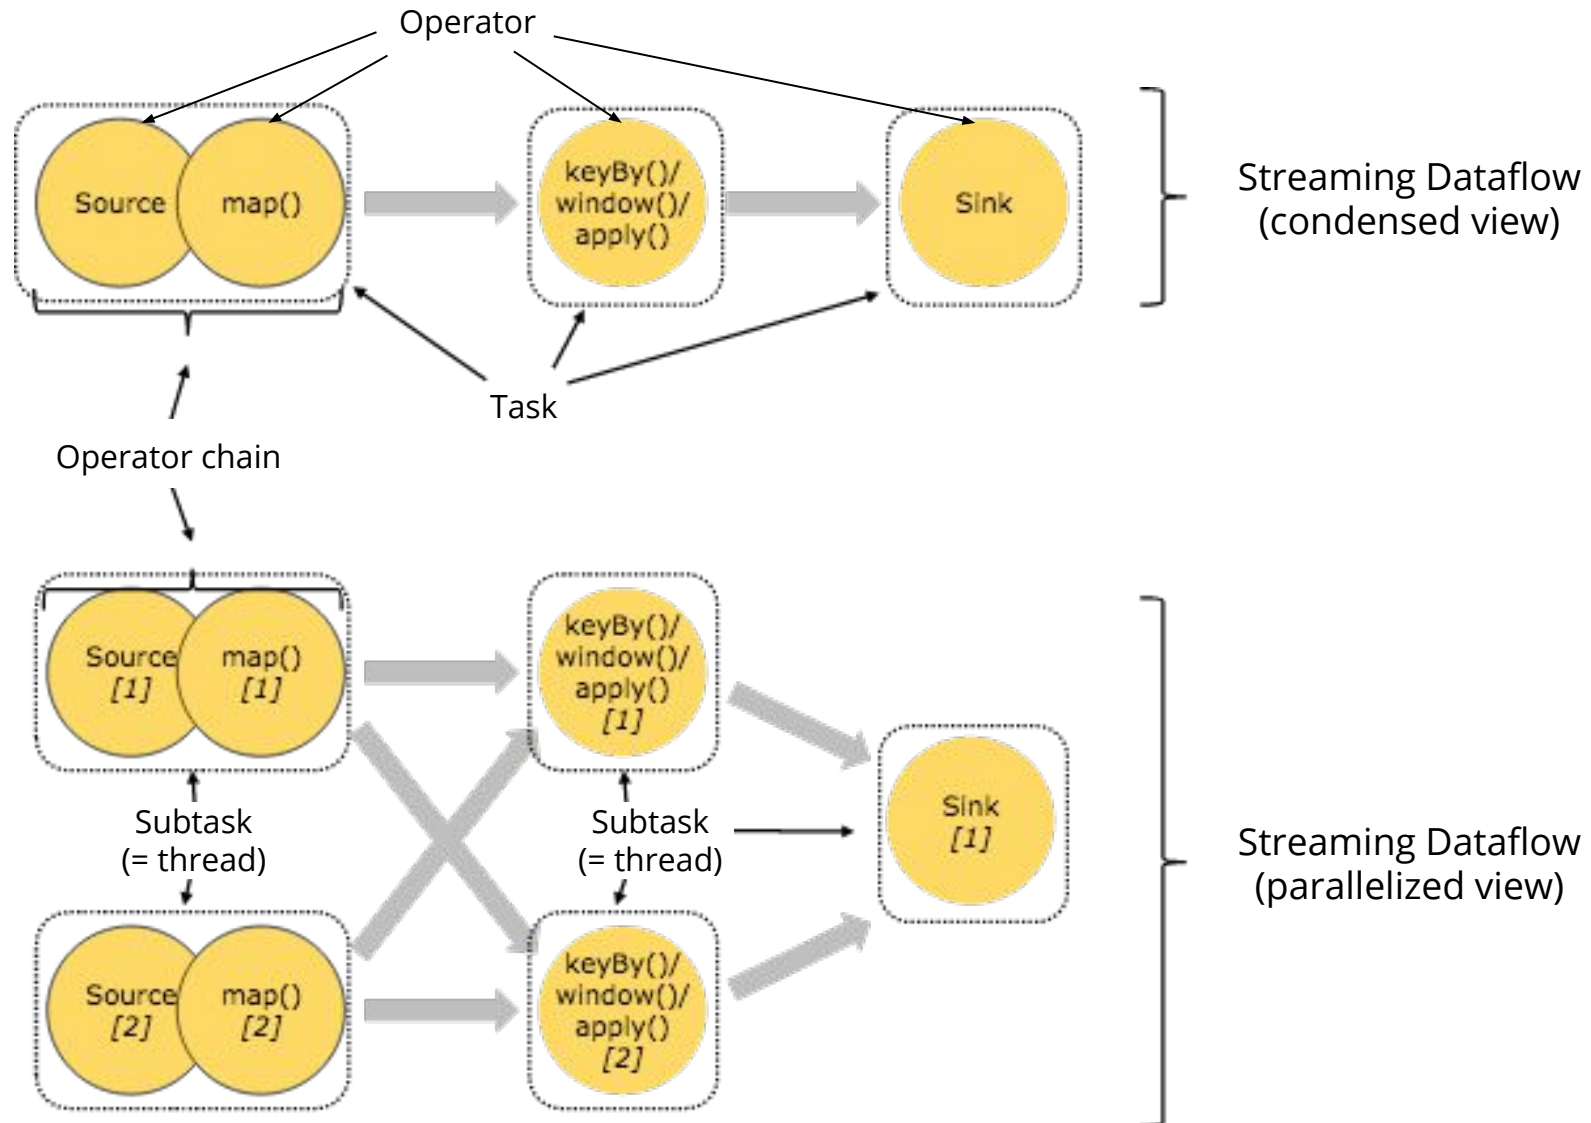
} Transformation

```
stats.addSink(new BucketingSink<>(path));
```

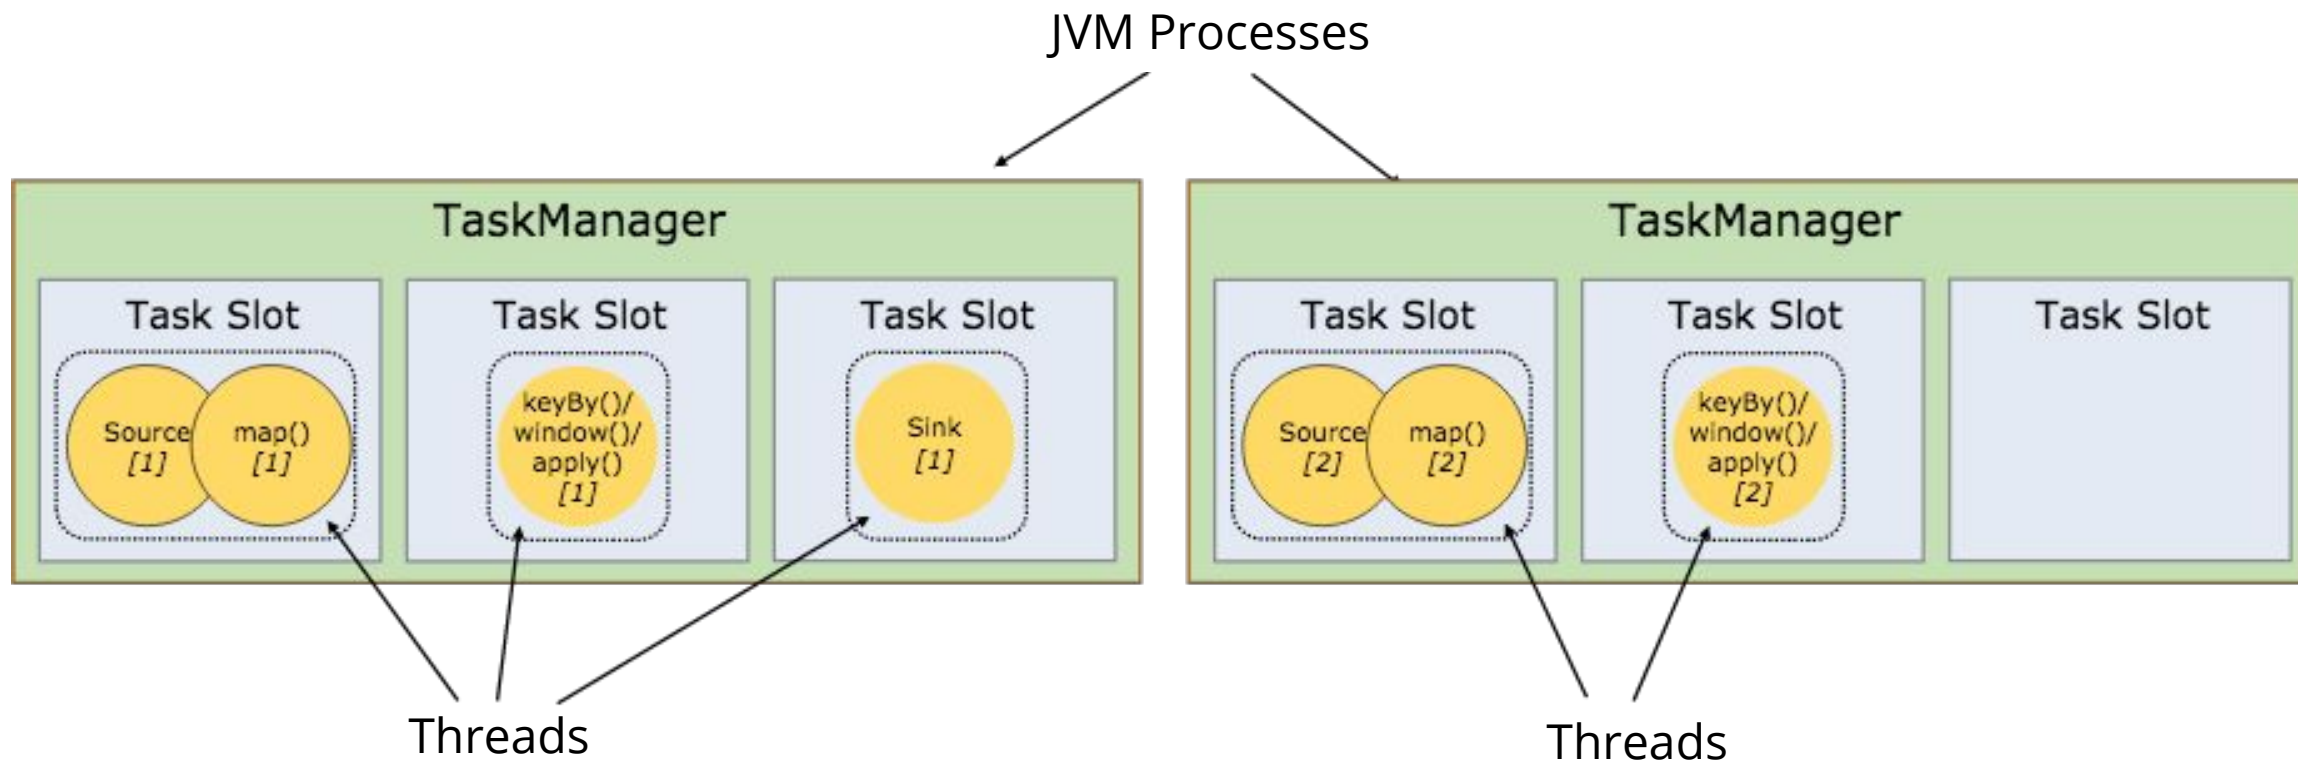
} Sink



Job Components & Operator chaining

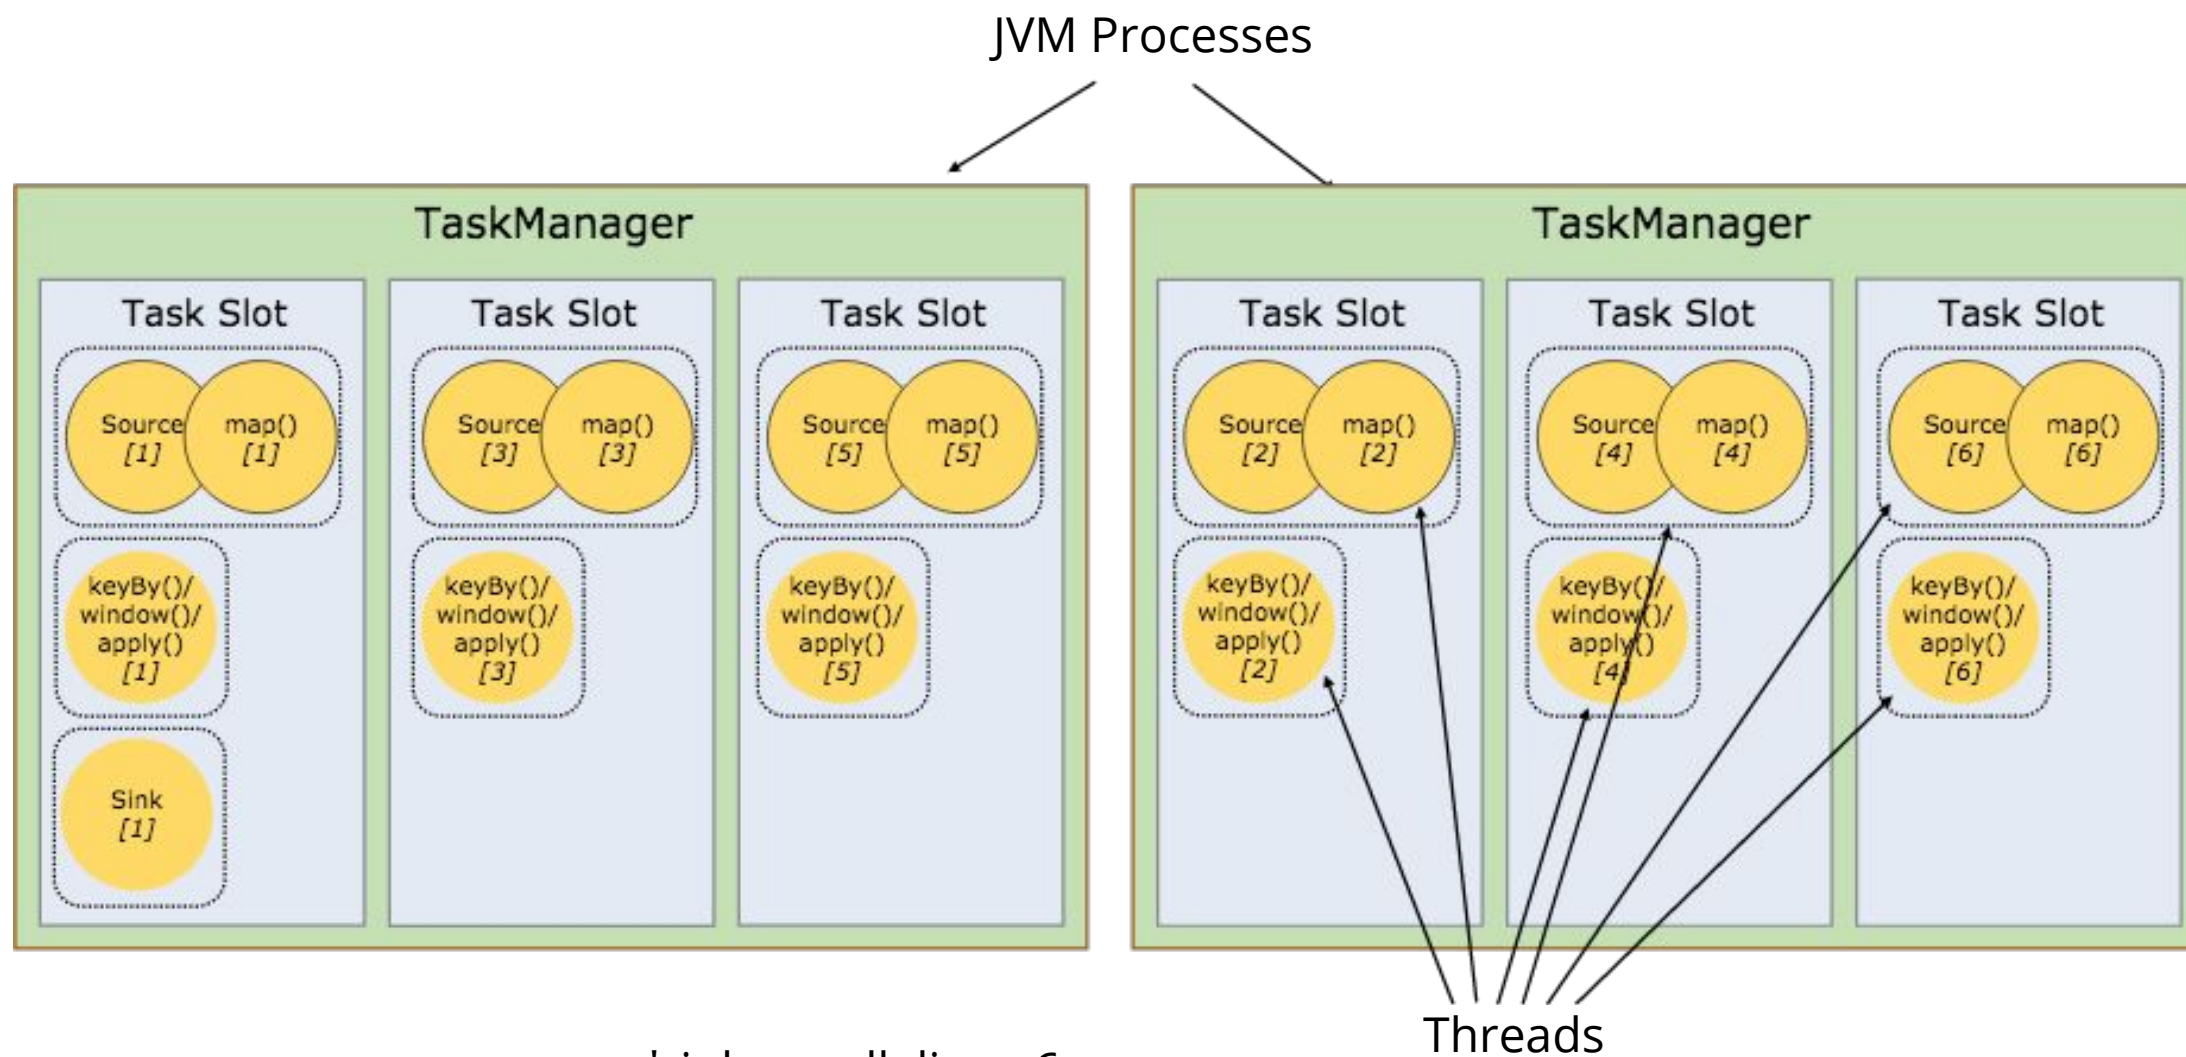


Task Deployment: Slots



* job parallelism: 2

Task Deployment: Slot Sharing



* job parallelism: 6

Fault Tolerance Guarantees

What happens if a worker goes down?

Flink supports different levels of guarantee for failure recovery:

Exactly once

- Each event affects the managed state exactly once.
- **Note:** This does *not* mean that events are processed exactly once!

At least once

- Each event affects the declared state of a program at least once

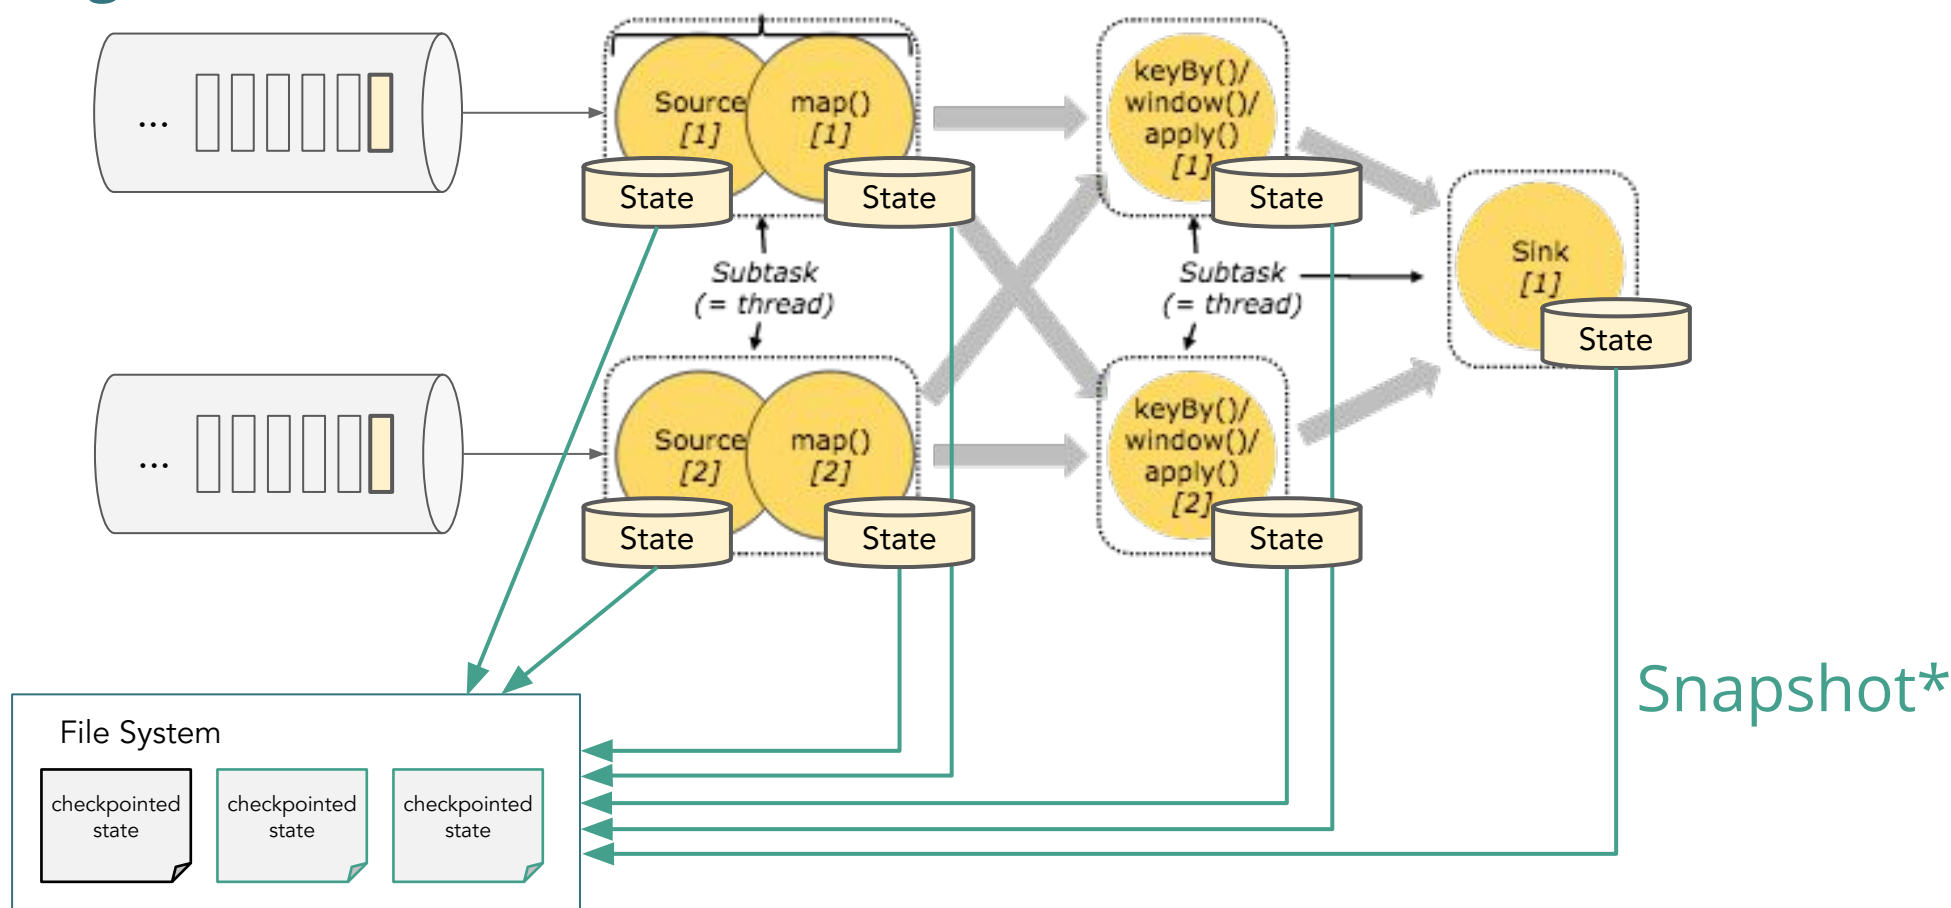
Deactivated / None / At most once

- All state is lost in case of a failure



Job Lifecycle & Fault Tolerance

While running

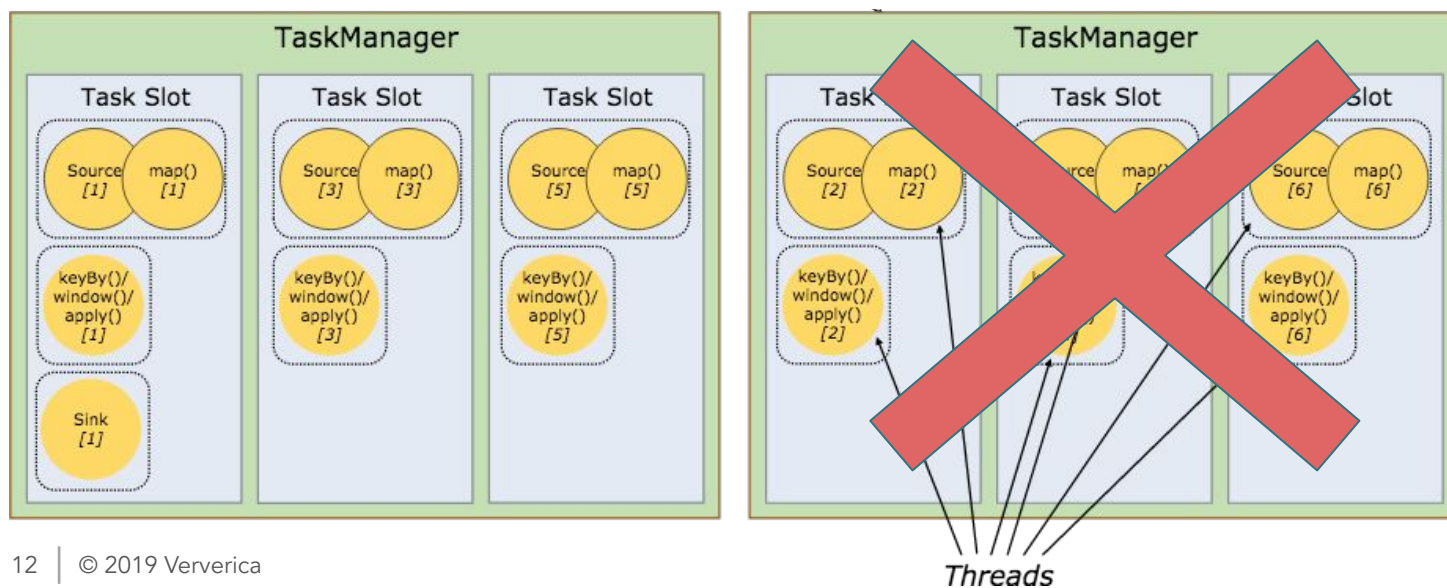
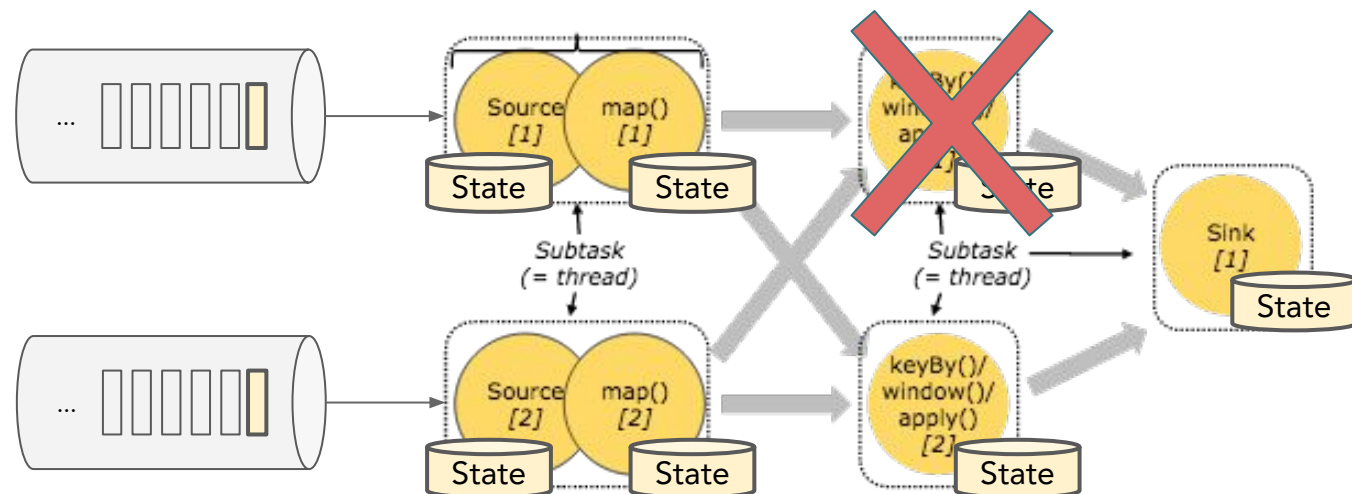


* Asynchronous Barrier Snapshotting
→ "checkpoint barriers" flow with data

Job Lifecycle & Fault Tolerance

On Failure

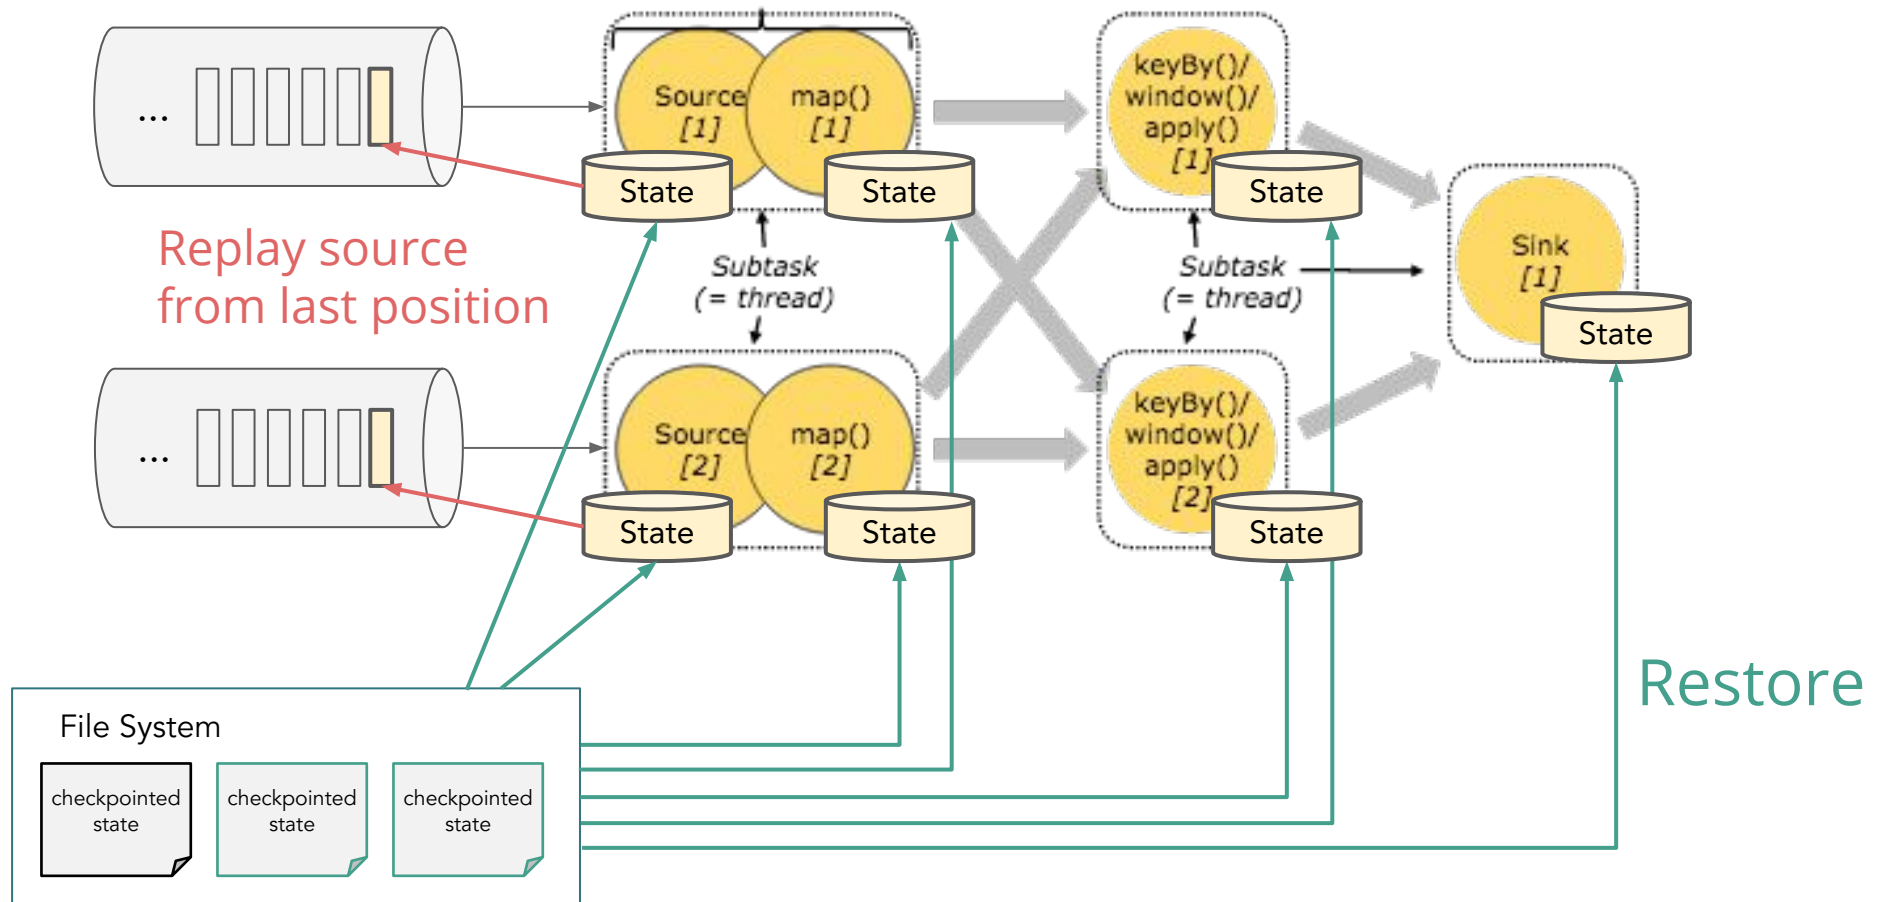
- TaskManager fails
⇒ cancel job on all TMs,
⇒ restart
(may require new TMs/slots)



- Operator code fails
⇒ cancel job on all TMs,
⇒ restart

Job Lifecycle & Fault Tolerance

After restart



Exercises

Exercises

Getting Started

1. Preparations

- a. `git clone https://github.com/ververica/flink-training-troubleshooting`
- b. `git pull origin master`
- c. Import project into IDE

2. Run the TroubledStreamingJob locally by executing TroubledStreamingJobRunner under test/

Note: The job is meant to restart periodically.

3. Run TroubledStreamingJob on Ververica Platform (see next slide)



Exercises

Getting Started

1. Login into Ververica Platform 2.0 (**pre-release!**)
 - `http://<your-training-ip>/` (Credentials provided)
2. Upload JAR
 - `mvn clean package`
 - Upload `target/flink-training-troubleshooting-0.1.jar` via ApplicationManager's Artifacts tab
3. Verify "Troubled Flink Job" Deployment in Ververica Platform
 - Upgrade Strategy: **STATELESS**, Restore Strategy: **NONE**, Parallelism: **4**
 - Jobmanager/Taskmanager CPU: **1/1 (defaults)**
 - Jobmanager/Taskmanager Memory: **2G/2G (defaults)**
4. Adapt Jar URI in Deployment in Application Manager if you change the name
5. Start "Troubled Flink Job" Deployment in Ververica Platform



Getting Started Walk-Through

Exercises

Debugging Frequent Job Failures

Exercise 1

As we just saw, `TroubledStreamingJob` is restarting frequently. Locate the problem and deploy a fix for it to Ververica Platform.

Note: Classes/Methods annotated with `@DoNotTouchThis` should not be modified during the training.





ververica

nico@ververica.com

www.ververica.com

[@VervericaData](https://twitter.com/VervericaData)