

Metrics & Monitoring

Nico Kruber, Solutions Architect & Apache Flink committer



Metrics & Monitoring

Agenda

- Flink's Metrics System *"How"*
 - Metrics
 - MetricsReporter
- Key Metrics for Continuous Monitoring *"What"*
 - Health
 - Throughput & Progress
 - Latency
- Key Metrics for Troubleshooting *"What else"*

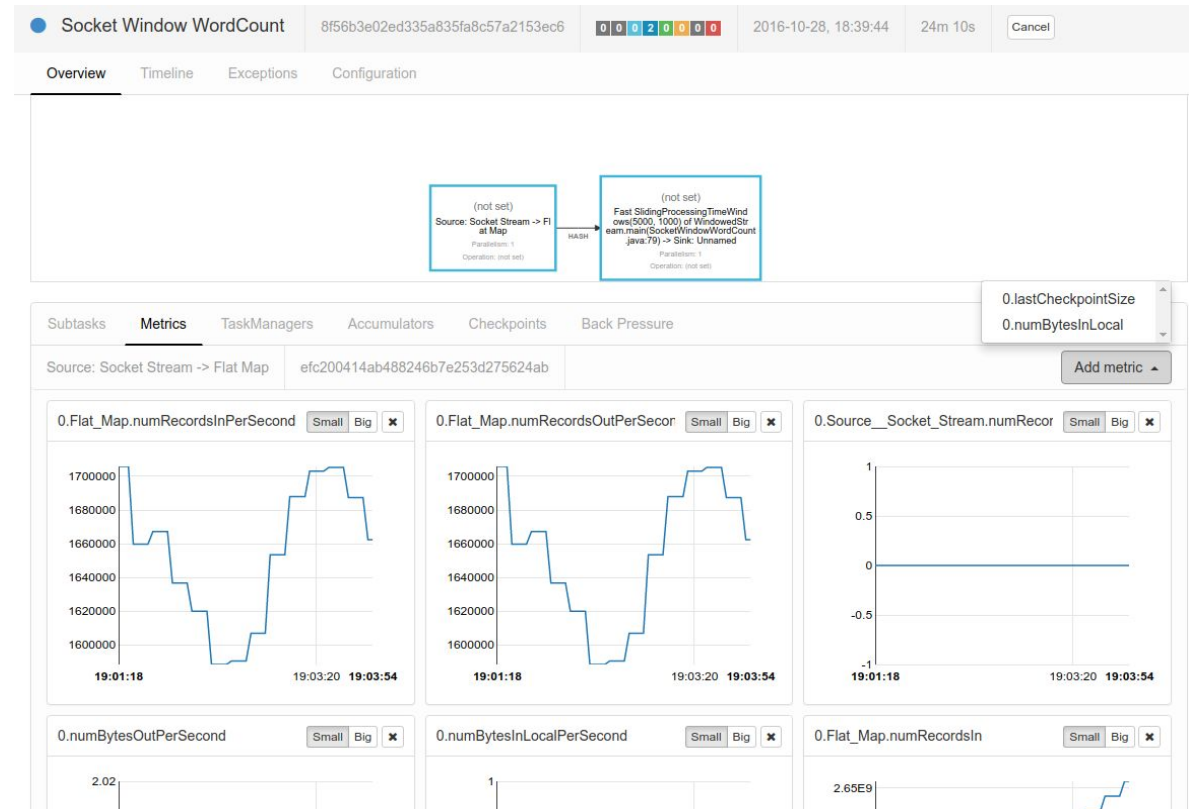


Flink's Metrics System



Metrics

- <identifier, measurement>
- Types
 - Counter
 - Meter (rate)
 - Histogram
 - Gauge (arbitrary value)



Example

```
public static class MyMap extends RichMapFunction<String, String> {  
    private Counter count;  
  
    @Override  
    public void open(Configuration config) {  
        count = getRuntimeContext()  
            .getMetricGroup()  
            .counter("numRecordsIn");  
    }  
  
    @Override  
    public String map(String input) {  
        count.inc();  
        // return something  
    }  
}
```



Metrics

Scopes

- metrics scope to different levels of a Flink deployment
- the keys to attach to metrics in a certain scope can be configured
 - `metrics.scope.jm: <host>.jobmanager`
 - `metrics.scope.task:`
`<host>.taskmanager.<tm_id>.<job_name>.<task_name>.<subtask_index>`
- Checkout <https://ci.apache.org/projects/flink/flink-docs-release-1.9/monitoring/metrics.html#scope> for details



Accessing Metrics

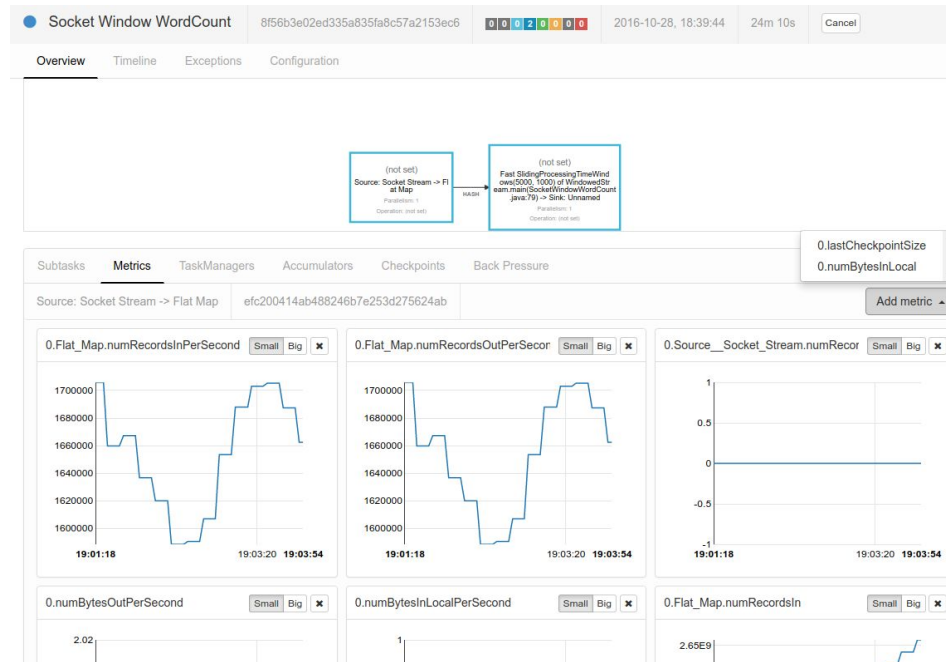
- WebUI → TaskMetrics
- REST API
- **MetricsReporters**

/jobs/<id>/metrics

/jobs/<id>/checkpoints

/jobs/<id>/vertices/<id>/metrics?get=0.numRecordsOutPerSecond

/taskmanagers/<id>/metrics?get=<metric>



Accessing Metrics

Metrics Reporters

- Datadog
- Ganglia
- Graphite
- JMX



- Prometheus
- StatsD
- SLF4J
- InfluxDB



Or write your own...



Accessing Metrics

A Simple Log4jReporter

```
public static class Log4JReporter implements MetricReporter, Scheduled {  
    private static final Logger LOG = LoggerFactory.getLogger(Log4jReporter.class);  
  
    private final Map<Counter, String> counters = new ConcurrentHashMap<>();  
  
    public void notifyOfAddedMetric(Metric metric, String metricName, MetricGroup group) {  
        if (metric instanceof Counter) {  
            counters.put((Counter) metric, group.getMetricIdentifier(metricName));  
        }  
    }  
  
    public void notifyOfRemovedMetric(Metric metric, String metricName, MetricGroup group) {  
        if (metric instanceof Counter) {  
            counters.remove(metric);  
        }  
    }  
  
    public void report() {  
        for (Map.Entry<Counter, String> metric : counters.entrySet()) {  
            LOG.info(metric.getValue() + ": " + metric.getKey());  
        }  
    }  
}
```



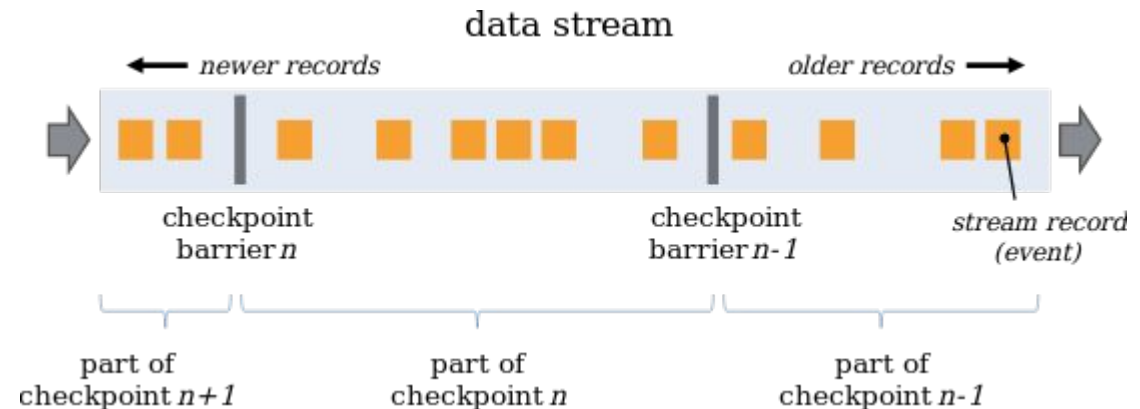
Key Metrics for Continuous Monitoring



Key Metrics

General Health

- Is “RUNNING”?
 - uptime
 - fullRestarts
- Checkpointing Consistently?
 - numberOfCompletedCheckpoints
 - numberOfFailedCheckpoints
 - lastCheckpointSize



Key Metrics

Throughput & Progress

- Task & Operator Level Throughput
 - `numRecords(In|Out)PerSecond`
 - `numRecords(In|Out)`
- Progress & Event-Time Lag
 - `currentOutputWatermark`
- Keeping Up
 - (Kafka) `records-lag-max`
 - (Kinesis) `millisBehindLatest`



Key Metrics

Latency

- Add timestamp to events at multiple stages, e.g.
 - event creation
 - ingestion
 - publishing
- custom metrics for reporting
 - these timestamps, or
 - the deviation from an ideal timestamp



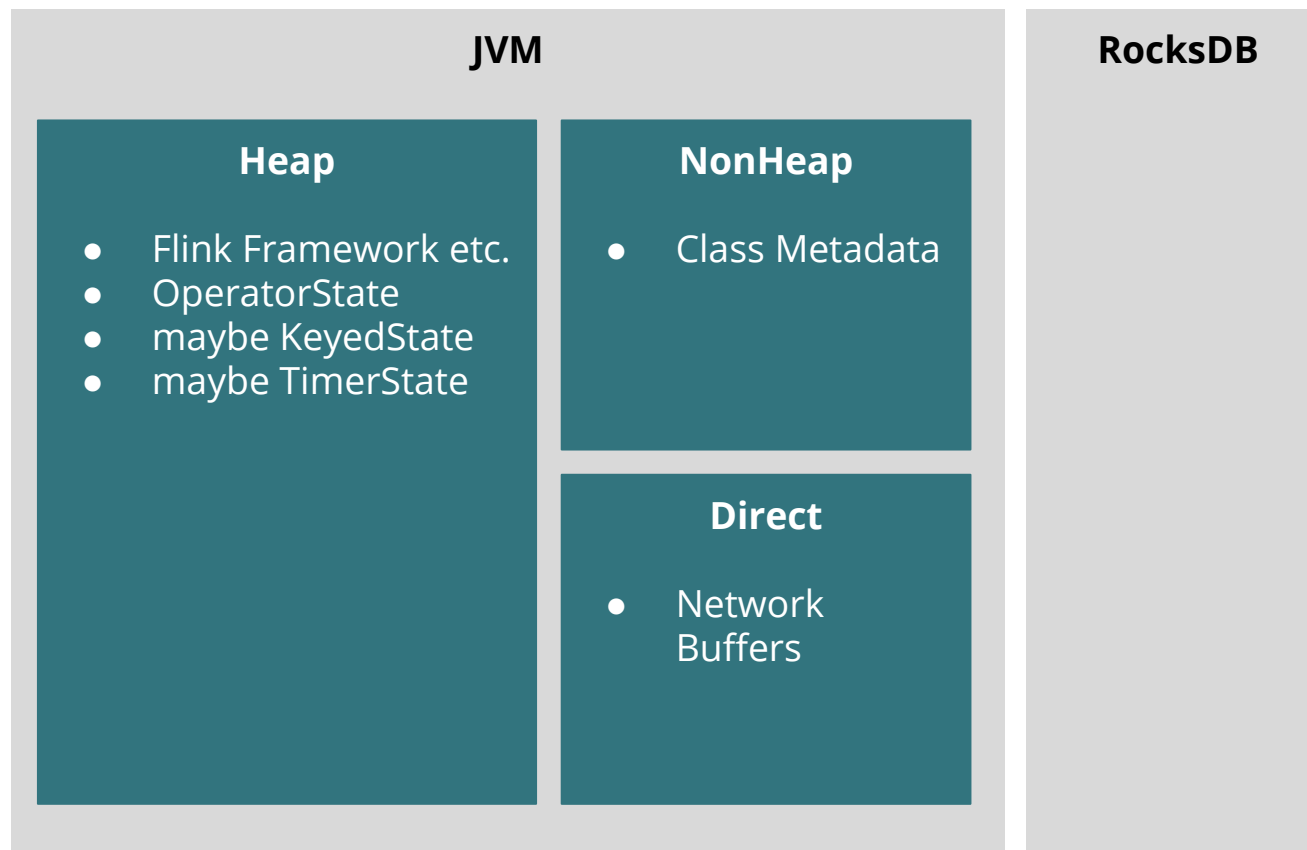
Key Metrics for Troubleshooting



JVM Metrics


Memory

- `Status.JVM.Memory.`
 - `NonHeap.Committed`
 - `Heap.Used`
 - `Heap.Committed`
 - `Direct.MemoryUsed`
 - `Mapped.MemoryUsed`
 - `G1 Young Generation.Time`
 - `G1 Old Generation.Time`



JVM Metrics

CPU

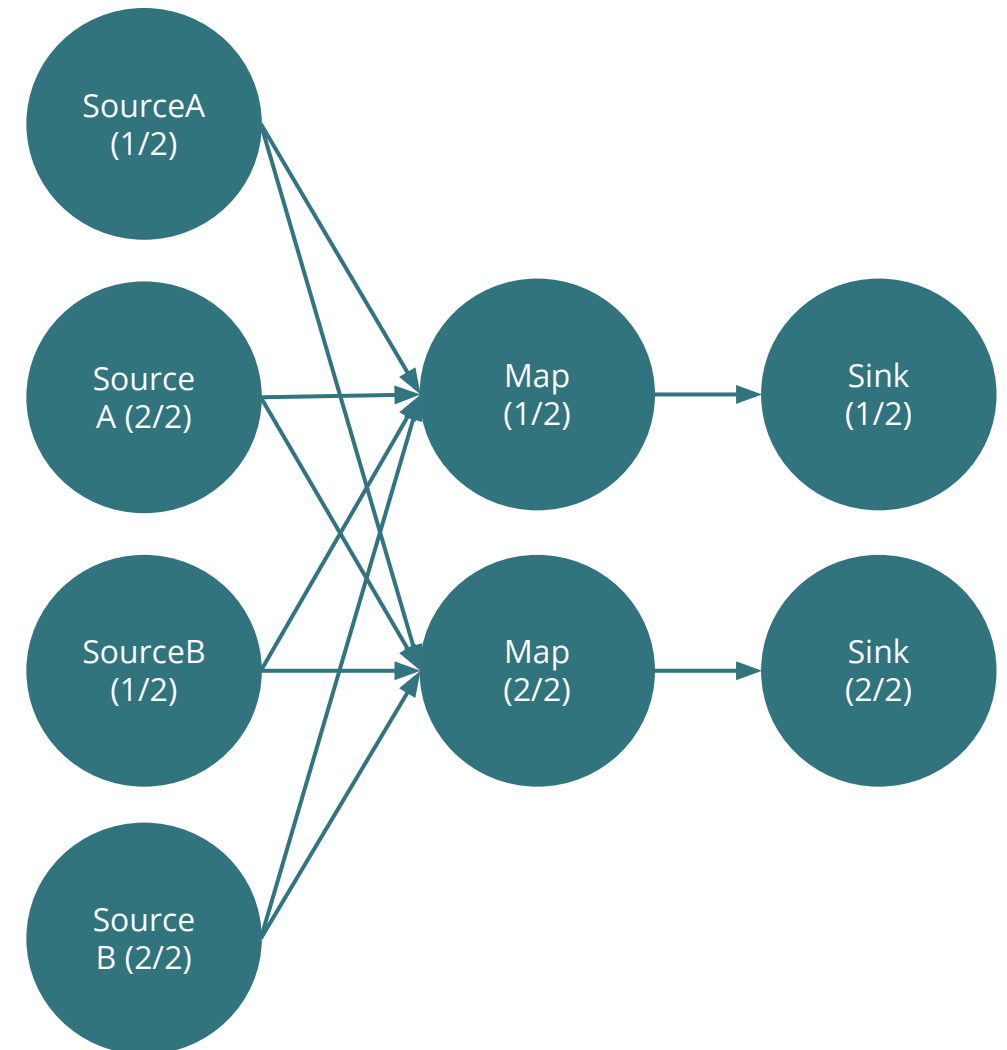
- Metrics
 - `Status.JVM.CPU.Load`
 - `Status.JVM.CPU.Time`
 - Leave some slack for catch-up scenarios (& RocksDB)
- Note:** 0.021 = 100% load for a Taskmanager container with 1 CPU on a 48 core machine.
- 



Troubleshooting Latency

Latency Tracking

- For each operator-subtask a latency histogram is exposed
- Enabled via `metrics.latency.interval`
- Scoped to job
- `latency.source_id.<source_id>`
 `.operator_id.<operator_id>`
 `.operator_subtask_index`
 `.<subtask_index>`

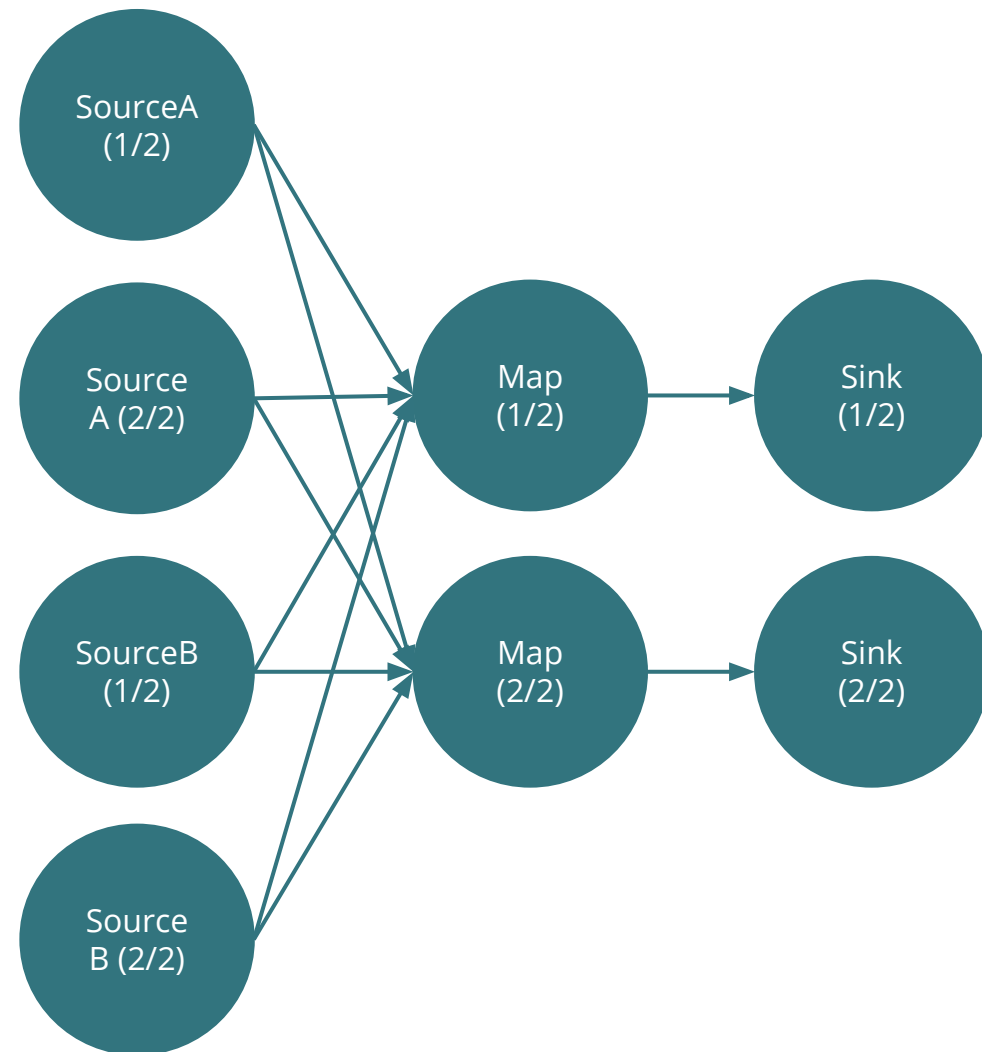


Troubleshooting Latency

Latency Tracking

metrics.latency.granularity: single

- Per Subtask
 - Latency histogram for both sources
- Overall
 - $4 (P * \#Operators)$

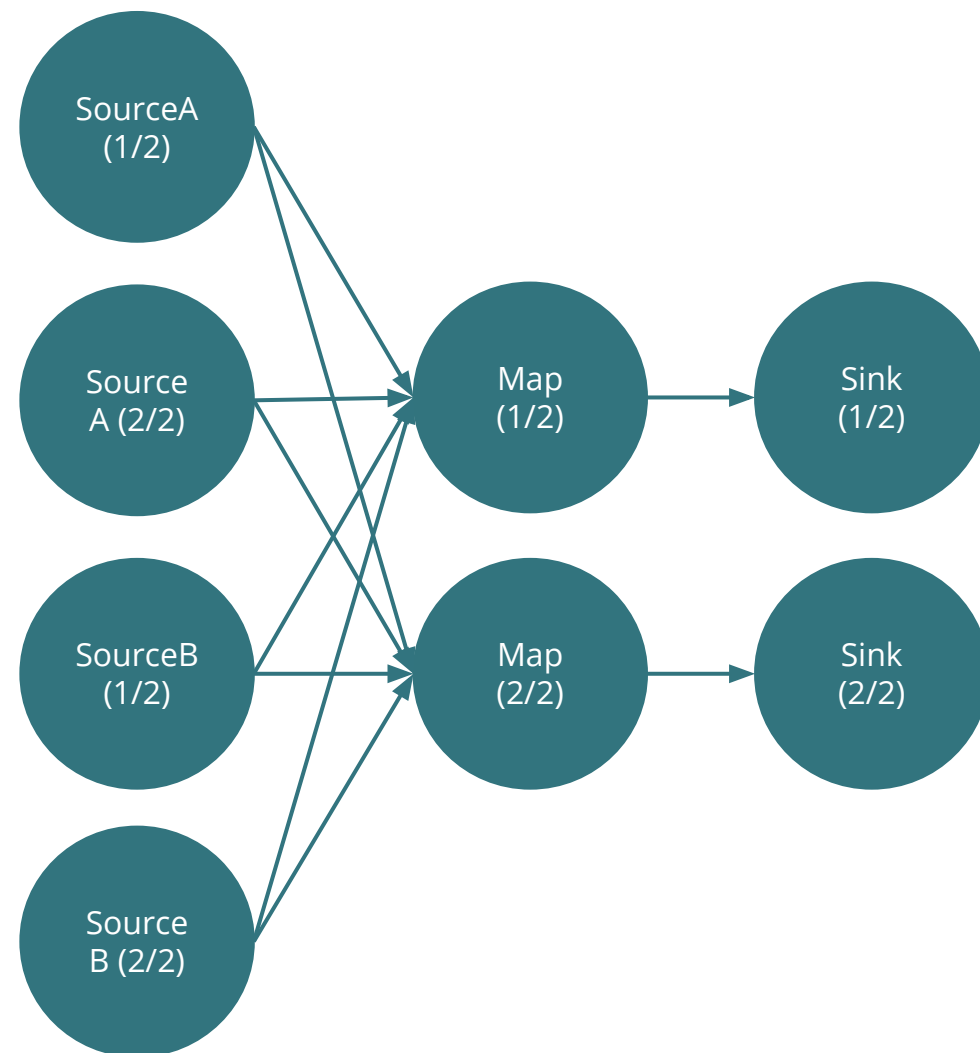


Troubleshooting Latency

Latency Tracking

metrics.latency.granularity: operator

- Per Subtask
 - Latency histogram for Source A
 - Latency histogram for Source B
- Overall
 - 8 histograms ($P * \#Sources * \#Operators$)

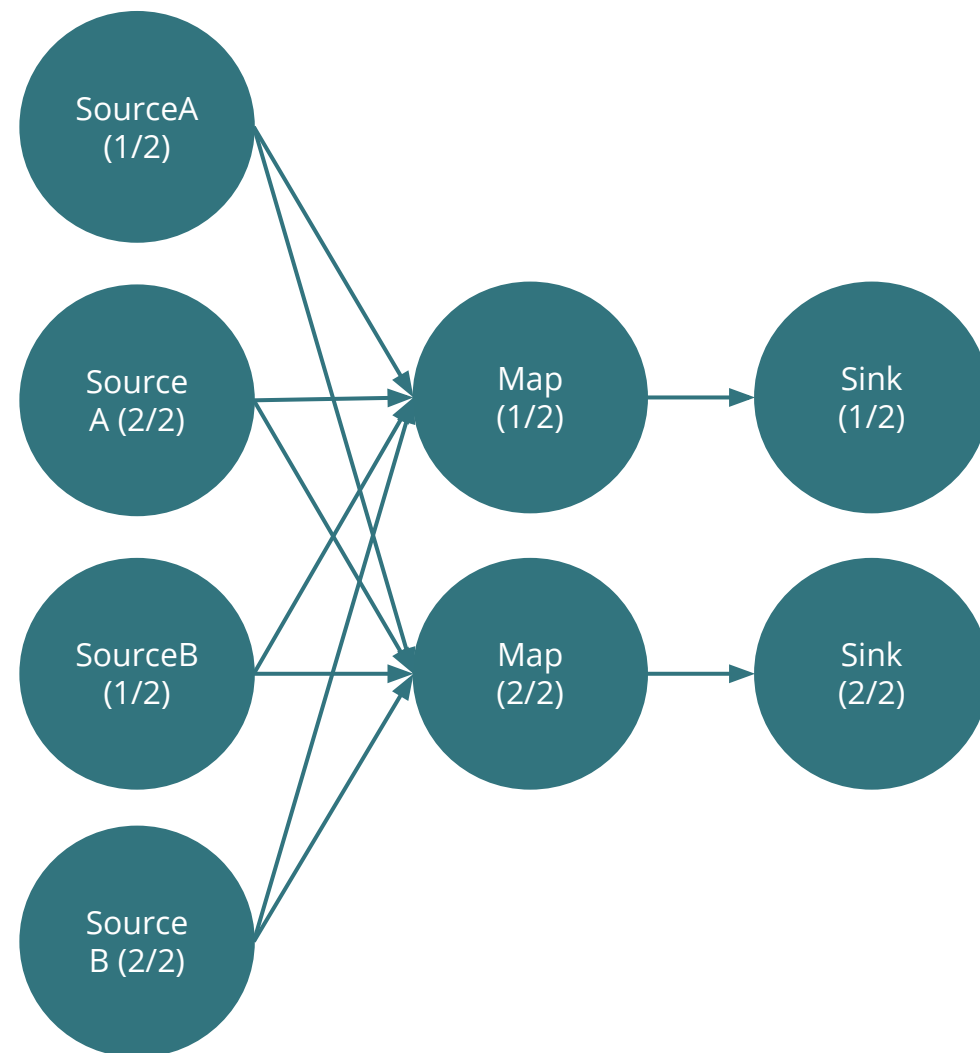


Troubleshooting Latency

Latency Tracking

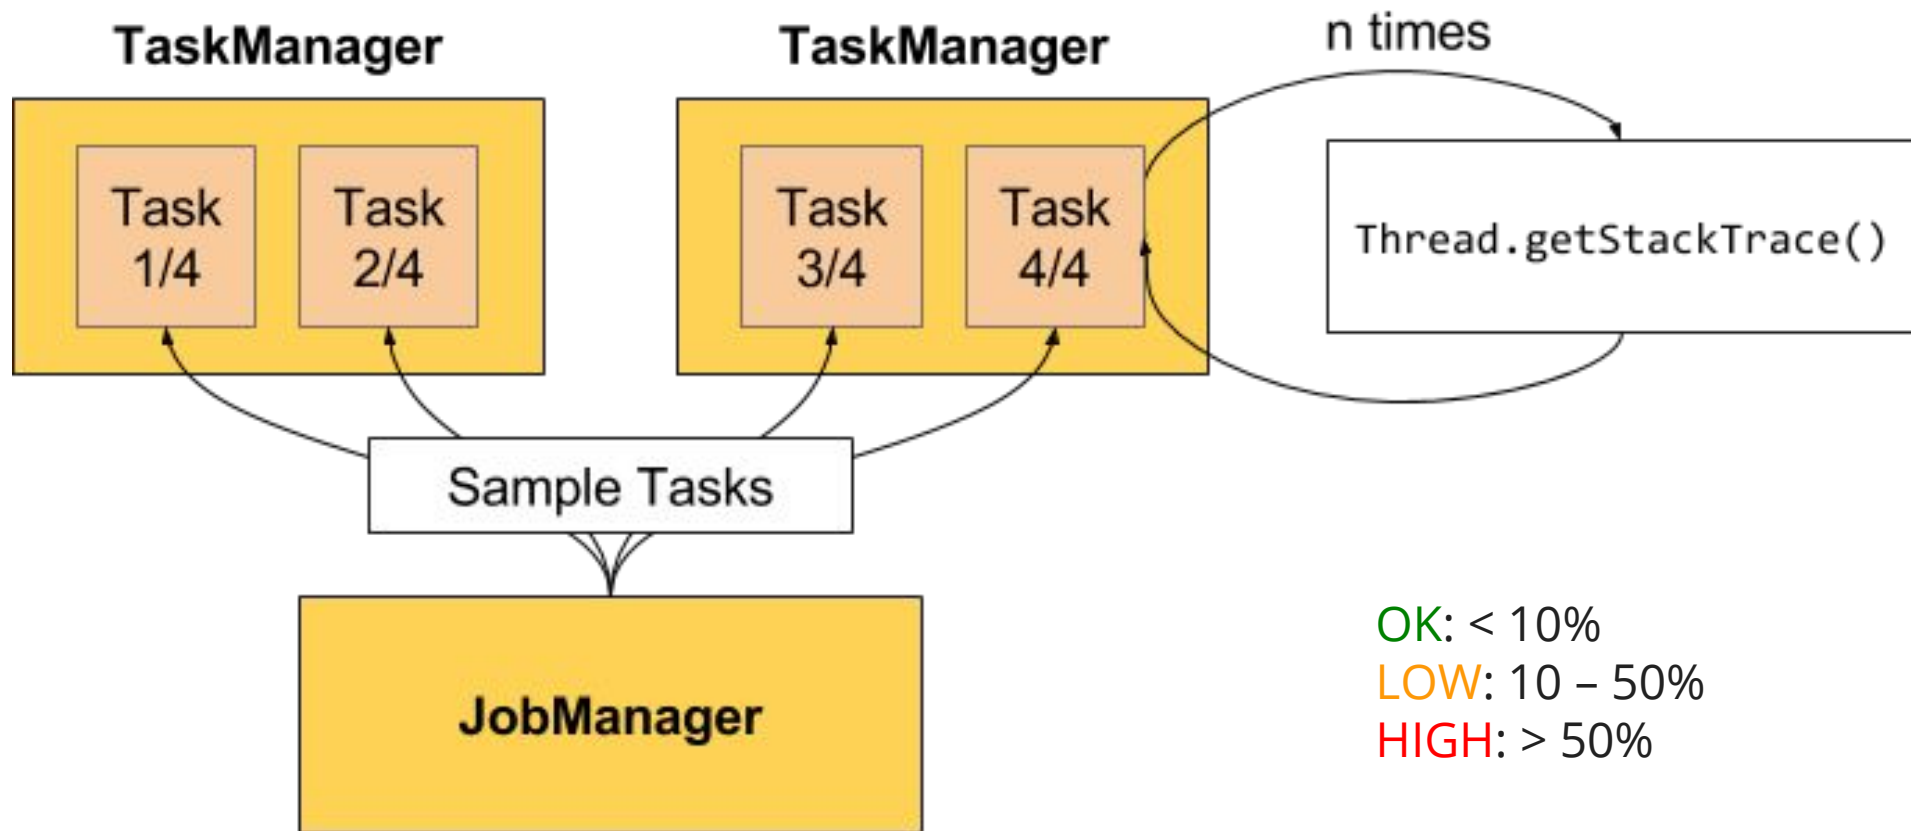
metrics.latency.granularity: subtask

- Per Subtask
 - Latency histogram for Source A (1/2)
 - Latency histogram for Source A (2/2)
 - Latency histogram for Source B (1/2)
 - Latency histogram for Source B (2/2)
- Overall
 - 16 histogram ($P^2 * \text{\#Sources} * \text{\#Operators}$)



Troubleshooting Backpressure

Backpressure Monitor



Measurement: Sampling in progress...

Back Pressure Status: -

SubTask

Ratio

Status



No Data

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Tasks
Sink: Print to Std. Out	RUNNING	184 MB	184	0 B	0	8	8
Flat Map	RUNNING	184 MB	184	184 MB	184	8	8
Map	RUNNING	184 MB	184	184 MB	184	8	8
Source: Custom Source	RUNNING	0 B	0	184 MB	184	4	4



Detail

SubTasks

TaskManagers

Watermarks

Accumulators

BackPressure

Metrics

Measurement: 17s ago

Back Pressure Status: OK

SubTask	Ratio	Status
1	0.01	OK
2	0	OK
3	0	OK
>	0	OK
5	0	OK
6	0.01	OK
7	0	OK
8	0	OK

...

Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Parallelism	Tasks
Sink: Print to Std. Out	RUNNING	1.71 GB	1,748	0 B	0	8	8
Flat Map	RUNNING	1.71 GB	1,748	1.71 GB	1,748	8	8
Map	RUNNING	1.71 GB	1,748	1.71 GB	1,748	8	8
Source: Custom Source	RUNNING	0 B	0	1.71 GB	1,748	4	4



Detail

SubTasks

TaskManagers

Watermarks

Accumulators

BackPressure

Metrics

Measurement: 1m 8s ago

Back Pressure Status: HIGH

SubTask	Ratio	Status
1	1	HIGH
2	1	HIGH
3	1	HIGH
>	1	HIGH
5	0.97	HIGH
6	1	HIGH
7	1	HIGH
8	1	HIGH

...

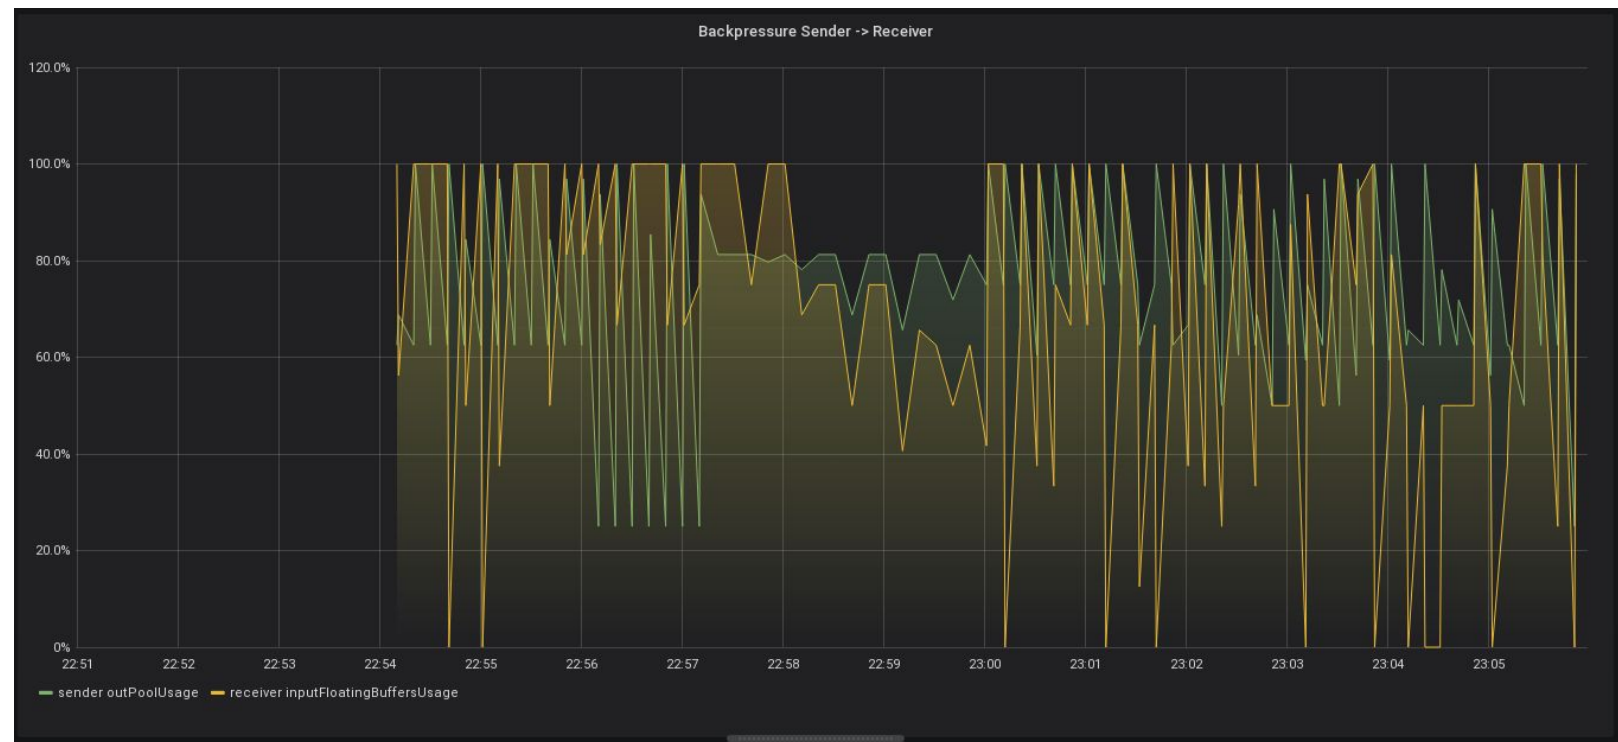
Name	Status	Bytes Received	Records Received	Bytes Sent	Records Sent	Tasks
Sink: Print to Std. Out	RUNNING	0 B	0	0 B	0	8
Flat Map	RUNNING	2.73 GB	2,792	0 B	0	8
Map	RUNNING	2.75 GB	2,800	2.73 GB	2,800	8
Source: Custom Source	RUNNING	0 B	0	2.75 GB	2,820	4



Troubleshooting Backpressure

Metrics

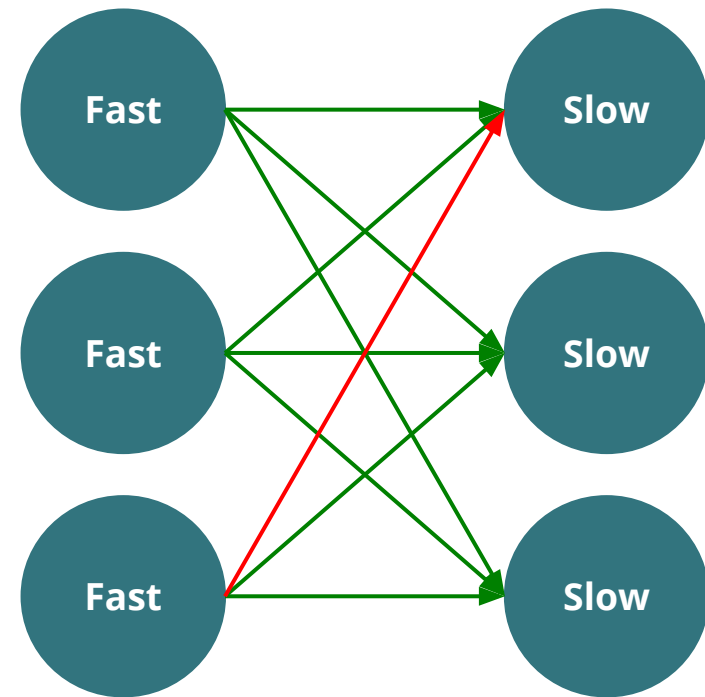
- sender's `outPoolUsage` vs. receiver's `inputFloatingBuffersUsage`
- monitored consistently
- may be used to identify the cause of backpressure
 - identify source operator
 - network latencies (credit-based flow control)
 - asymmetric backpressure
 - other resource bottleneck



Troubleshooting Backpressure

Asymmetric Backpressure

- situation where backpressure only occurs in one channel
- hard to detect, but can lead to checkpoint timeouts
- Metrics
 - `inputFloatingBuffersUsage`,
`inputExclusiveBuffersUsage`
 - `outPoolUsage`





ververica

nico@ververica.com

www.ververica.com

[@VervericaData](https://twitter.com/VervericaData)