





Object Reuse & Serialization (2)

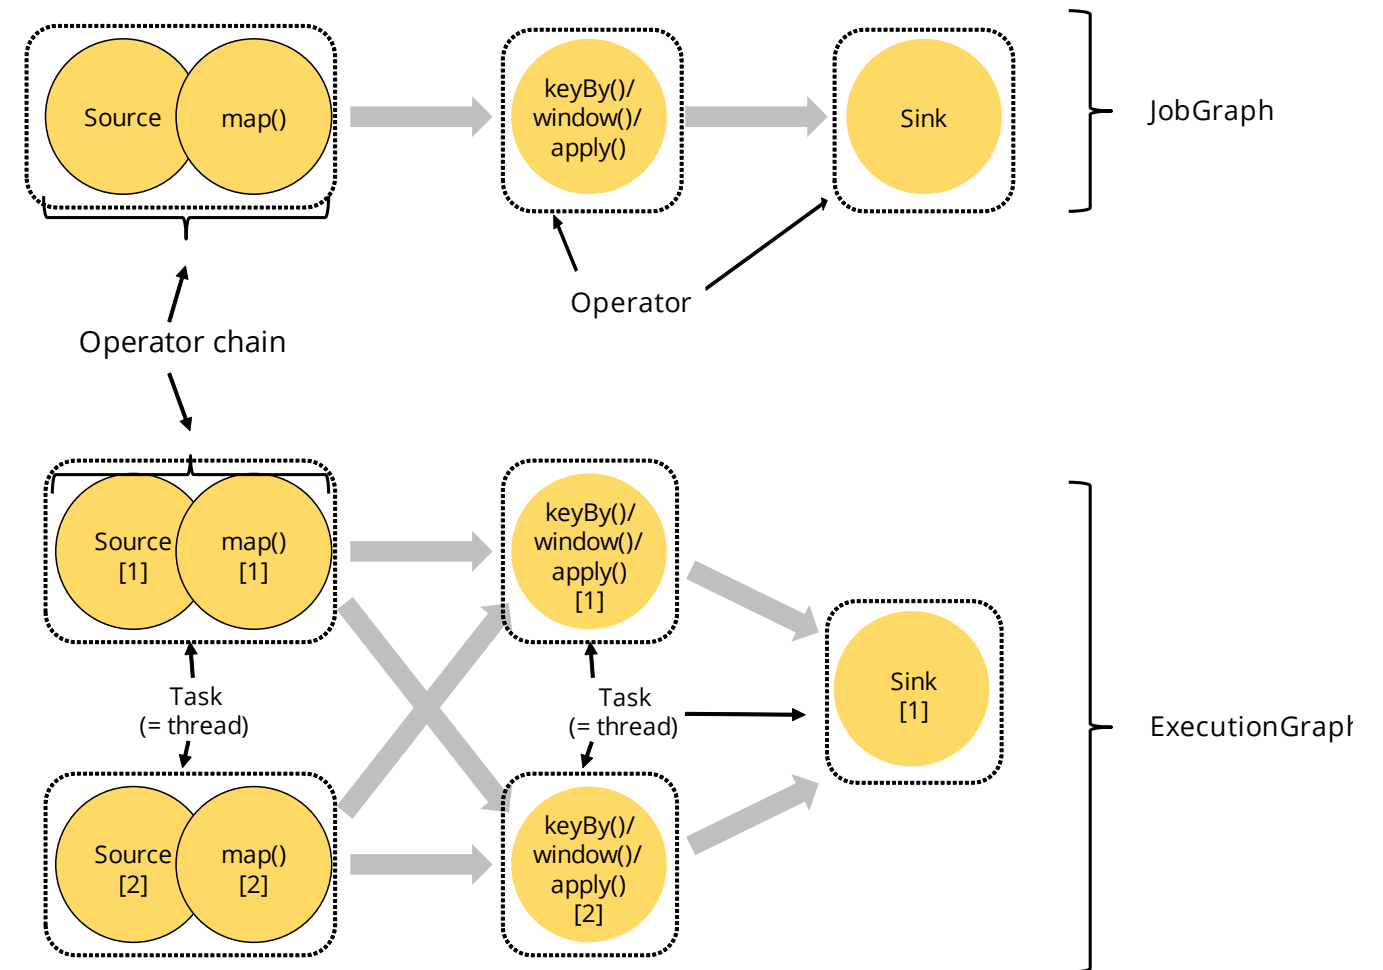
Apache Flink[®] Tuning & Troubleshooting Training

Recall: Data exchange strategies



Recall: Operator chaining

- Chaining:
 - Sender's `out.collect()` directly calls into receiver
 - Records are passed along with a defensive copy
 - but without any of the lower layers of Flink (de/ser*, network stack)
 - Only possible during **forward** data exchange



Object Reuse

- disabled by default
 - `ExecutionConfig#enableObjectReuse()`
- when enabled
 - streaming: same object is used throughout operator chains
 - batch: more complicated re-use patterns possible <https://ci.apache.org/projects/flink/flink-docs-stable/dev/batch/index.html#object-reuse-enabled>

Object Reuse in Streaming

Effects

Without a defensive between two operators

- Second operator could change a value the first one is holding
 - stored in field in first operator code
 - stored in (heap-based!) state back-end
 - first operator accesses value after `out.collect()`
- Value storage can be implicit
 - Window state
 - Context key
- Applies to method parameters, Iterables,...

Object Reuse in Streaming

Restrictions

- It is **not safe** to remember input objects across function calls.
- You **must not** modify input objects.
- You *may* modify an output object and emit it again (following the rules above).



Immutable Types (Streaming)

- Serialization via `PojoSerializer` requires getters **and** **setters** (or public fields).
 - Otherwise falls back to Kryo
- Non-guarded access vs. slow serialization
- Can work around by providing custom serializers.
- Can distinguish between
 - Wire serialization
 - State serialization



Custom Serialization

```
@TypeInfo(MyTupleTypeInfoFactory.class)
public class MyTuple<T0, T1> {
    public T0 myfield0;
    public T1 myfield1;
}

public class MyTupleTypeInfoFactory extends TypeInfoFactory<MyTuple> {

    @Override
    public TypeInformation<MyTuple> createTypeInfo(
        Type t, Map<String, TypeInformation<?>> genericParameters) {
        return new MyTupleTypeInfo(genericParameters.get("T0"), genericParameters.get("T1"));
    }
}
```



Exercises

Exercise 5

- Run `ObjectReuseJob` and observe the provided graphs in Grafana.
- Improve throughput by enabling object reuse
 - adapt Job arguments in deployment: `mainArgs: '--objectReuse true'`.
- What are the effects? **Try to fix incorrect results.**

Bonus Task 5.1

Make sure that data types are immutable. What are the effects?

Bonus Task 5.2

Use a custom serializer to avoid serialization via Kryo. Compare object-reuse `enabled` and `disabled`.





ververica